



Computer
Science

CSC380: Principles of Data Science

Data Analysis, Collection, and Visualization 1

Xinchen Yu

Announcements

2

No in-person lecture on Feb 8 (Thurs) and Feb 13 (Tues).

Recordings of Data Summarization will be uploaded to D2L.

HW3 out this Thursday, due next Friday (individual work).

- Data Collection and Sampling
- Data Visualization
- Data Summarization

Not understanding how data are collected is one of the top reasons behind bad data science...

How Bad Data Is
Undermining Big Data
Analytics

Forbes

How to be a bad data scientist!



Pascal Potvin Feb 27, 2018

**8 telltale signs of
a bad data scientist**

InfoWorld

**If Your Data Is Bad, Your
Machine Learning Tools
Are Useless**

by Thomas C. Redman

...we will not do data collection or experimental design, but should be familiar with the basics

1. Plan research design
2. Collect data (essentially, sampling) **Will focus on these**
3. Visualize and summarize the data (plots and summary stats)
4. Make inferences from data (i.e., estimate stuff, test hypotheses, ...)
5. Interpret results

Randomized Control. Researcher controls treatment among groups. Used to assess *causal* relationships. Stronger than correlational study but difficult to conduct. (e.g., clinical trials)

Observational. Collect data by “observing” passively. If there are treatments (i.e., vaccines), they are not under control of the researcher.

- **Natural Experiment.** Observe naturally-occurring phenomena. Approximates a controlled study, despite the researcher not having control of any groups. (e.g. Helena, Montana banned smoking ban in all public spaces for six months –before/after this ban)
- **Case Studies and Surveys.** Analysis based on previously-collected data. (e.g., Analysis of US census data, or US current population survey (CPS))

Recall: Explanatory, response variable

7

Example: Say we study the relationship between **Smoking** vs **Cancer**

Independent variable: variables that are manipulated or are changed by researchers and whose effects are measured and compared.

= **explanatory variable**

Dependent variable: the variable that depends on independent variable (or speculated to do so).


= **response variable**

Recall: Explanatory, response variable

Research question	Independent variable	Dependent variable(s)
Do tomatoes grow fastest under fluorescent, incandescent, or natural light?	<ul style="list-style-type: none">Type of light the tomato plant is grown under	<ul style="list-style-type: none">The rate of growth of the tomato plant
What is the effect of intermittent fasting on blood sugar levels?	<ul style="list-style-type: none">Presence or absence of intermittent fasting	<ul style="list-style-type: none">Blood sugar levels
Is medical marijuana effective for pain reduction in people with chronic pain?	<ul style="list-style-type: none">Presence or absence of medical marijuana use	<ul style="list-style-type: none">Frequency of painIntensity of pain
To what extent does remote working increase job satisfaction?	<ul style="list-style-type: none">Type of work environment (remote or in office)	<ul style="list-style-type: none">Job satisfaction self-reports


Causation vs. Correlation

Studies generally try to show *either* correlation (association) or causation, but they are not the same...




A variable that influences the *response* but is unaccounted for in data collection

Example: You are studying whether **birth order** affects **Down's Syndrome** in the child. You collect samples of children, their birth order, and cases of Down's syndrome.



- You went on to collecting the maternal age data.




So.. both maternal age and birth order is associated with Down's syndrome?

A variable that influences the *response* but is unaccounted for in data collection


Example: You are studying whether **birth order** affects **Down's Syndrome** in the child. You collect samples of children, their birth order, and cases of Down's syndrome.

Explanation: Maternal age (confounder) was not recorded. Two scenarios:

1. Higher maternal age is directly associated with Down's syndrome, regardless of birth order.
2. Maternal age directly assoc. with birth order (mother is older with later children), but not directly associated with Down's syndrome.



Stratified Sampling: Divide population into smaller groups. Previous example can divide population of children by maternal age at birth and collect data from each stratum




Approach


1. **Control** confounders: design treatments
2. **Randomize** the assignment of subjects to treatments (to eliminate bias due to systematic differences in categories)
3. **Replicate** experiment on many subjects, to obtain statistically meaningful results

Placebo Effect

Placebo Effect

The placebo effect is when a person has a response to a fake treatment.



 the tendency of any medication or treatment, even an inert or ineffective one, to exhibit results simply because the recipient believes that it will work.

1. Placebo Control: Subjects are randomly selected to receive either the vaccine or an injection of saline solution

2. Randomize: Stratified sampling with age strata: 12-15yrs, 16-55yrs, 55+yrs

3. Replicate: Experiment is repeated at multiple sites in several countries

Example: Pfizer COVID Phase 3 Vaccine Trials

17

The landmark phase 3 clinical trial enrolled **46,331** participants at **153** clinical trial sites around the world.

Trial Geography



Our trial sites are located in **Argentina, Brazil, Germany, Turkey, South Africa** and the **United States**.

Participant Diversity

Approximately **42%** of overall and **30%** of U.S. participants have diverse backgrounds.

Participants	Overall Study	U.S. Only
Asian	5%	6%
Black	10%	10%
Hispanic/Latinx	26%	13%
Native American	1.0%	1.3%

49.1% of participants are male and **50.9%** are female

Participant Age



Ages 12-15 2,260


Ages 16-17 754

Ages 18-55 25,427

Ages 56+ 17,879

Data Collection

Generally infeasible to collect data from entire *population*



Population Entire group that we want to draw conclusions about.

Can be defined in terms of location, age, income, etc.

Sample Specific group that we collect data from.

Necessity It is usually impractical or impossible to collect data from an entire population due to size or inaccessibility.


Cost-effectiveness There are fewer participant, laboratory, equipment, and researcher costs involved.

Manageability Storing data and running statistical analyses is easier on smaller datasets.

Population parameter A measure that describes *the whole population*.

Sample statistic A measure that describes the sample and reflects the population parameter.

Example We are studying student **political attitudes** and ask students to rate themselves on a scale: 1, very liberal, to 7, very conservative. The **population parameter** of interest is the average political leaning. The sample mean, say 3.2, is our **statistic**.



The *sampling error* is the difference between the population parameter and the sample statistic.

- Sampling errors are **normal**, but we want them to be low
- Samples are **random**, so sample statistics are estimates and thus subject to random noise
- **Sample bias** occurs when the sample is not representative of the population (for various reasons)

Sampling must be conducted properly, to avoid sample bias

Two primary types of sampling...

Probability Sampling Random selection allowing strong statistical inferences about the population


Non-Probability Sampling Based on convenience or other criteria to easily collect data (but no random sampling)




Simple random sample




Systematic sample



Stratified sample



Cluster sample



Simple Random Sample (SRS)

Each member of the population has the *same chance* of being selected (i.e., uniform over the population)

Example : American Community Survey (ACS)


Each year the US Census Bureau use *simple random sampling* to select individuals in the US. They follow those individuals for 1 year to draw conclusions about the US population as a whole.




Simple random sample




Systematic sample



Stratified sample



Cluster sample



Simple Random Sample (SRS)


Each member of the population has the *same chance* of being selected (i.e., uniform over the population)

- Most straightforward probability sampling method
- Impractical unless you have a complete list of every member of population


Simple random sample




Systematic sample



Stratified sample



Cluster sample



Systematic Sample

Select members of population at a regular interval, determined in advance


Example You own a grocery store and want to study customer satisfaction. You ask *every 20th customer* at checkout about their level of satisfaction.

Note We cannot itemize the whole population in this example, so SRS is not possible.


Simple random sample




Systematic sample



Stratified sample



Cluster sample



Systematic Sample

Select members of population at a regular interval, determined in advance


- Imitates SRS but is easier in practice
- **Do not** use when there can be a pattern. E.g., survey at the exit of a rollercoaster with N seats but with every N -th customer.

Alternative: use a Bernoulli(p) (e.g., $p=1/20$)


Simple random sample




Systematic sample



Stratified sample



Cluster sample



Stratified Sample


Divide population into *homogeneous* subpopulations (strata). Probability sample the strata.

Example We wish to solicit opinions of UA CS freshman by asking 100 of them, but they are about 14% women. SRS could easily fail to capture adequate number of women. We divide into men / women and perform SRS within each group.


Simple random sample



Systematic sample



Stratified sample



Cluster sample



Stratified Sample


Divide population into *homogeneous* subpopulations (strata). Probability sample the strata.

- Use when population is diverse and want to accurately capture characteristic of each group
- Ensures similar variance across subgroups
- Lowers overall variance in the population


Simple random sample




Systematic sample



Stratified sample



Cluster sample



Cluster Sample


Divide population into subgroups (clusters). Randomly select entire clusters.

Example We wish to study the average reading level of *all 7th graders in the city* (population). Create a list of all schools (clusters) then randomly select a subset of schools and test every student.


Simple random sample




Systematic sample



Stratified sample



Cluster sample




Cluster Sample

Divide population into subgroups (clusters). Randomly select entire clusters.

- This is *single-stage* cluster sampling
- *Multi-stage* avoids sampling every member of a group
- Related to stratified sampling, but groups are not homogeneous

Sampling Bias

Occurs if data are collected in a way that some members of the population have lower/higher probability of being sampled than others





Sometimes is unavoidable (e.g., not all members are equally accessible) but
(1) you should be aware of it
(2) must be corrected if possible at all

Example We conduct a poll by randomly calling numbers in a phone book. People that have less time are less likely to response. Called **non-response bias**.

Self-selection: Possible whenever members under study have control over whether to participate.

- E.g., online or phone-in poll—user can choose whether to initiate participation.

Exclusion: Excluding certain groups from the sample.

- E.g., longitudinal data collection for 2 years in a town where we exclude groups that move in or out.

Survivorship: Only *surviving* subjects are selected.

- E.g., studying improvements on the retention rate of CS majors by surveying CS graduates.

SRS is least prone to bias, but not always...

You want to study procrastination and social anxiety levels in undergraduate students at your university using a simple random sample. You assign a number to every student in the research participant database from 1 to 1500 and use a random number generator to select 120 numbers.

What is the cause of bias in this simple random sample?

SRS is least prone to bias, but not always...

You want to study procrastination and social anxiety levels in undergraduate students at your university using a simple random sample. You assign a number to every student in the research participant database from 1 to 1500 and use a random number generator to select 120 numbers.

Although you used a random sample, not every member of your target population –undergraduate students at your university – had a chance of being selected. Your sample misses anyone who did not sign up to be contacted about participating in research. This may bias your sample towards people who have less social anxiety and are more willing to participate in research.