



Computer  
Science

# CSC380: Principles of Data Science

## Linear Models 1

Xinchen Yu

Let  $m_c := \sum_{i=1}^m \mathbf{I}\{y^{(i)} = c\}$  be number of training examples in class  $c$  then,

$$\sum_{i=1}^m \log p(\mathcal{D}; \pi, \theta) = \sum_{c=1}^C m_c \log \pi_c + \sum_{c=1}^C \sum_{i: y^{(i)}=c} \sum_{d=1}^D \log p(x_d^{(i)}; \theta_{cd})$$

Log-likelihood function is concave in all parameters so...

1. Take derivatives with respect to  $\pi$  and  $\theta$  separately.
2. Set derivatives to zero and solve

$$\hat{\pi}_c = \frac{m_c}{m}$$

**Fraction of training  
examples from class  $c$**

$$\hat{\theta}_{cd} = \frac{m_{cd}}{m_c}$$

**Number of “heads” in  
training set from class  $c$**

$$m_{cd} = \sum_{i=1}^m \mathbf{I}\{y^{(i)} = c, x_d^{(i)} = 1\}$$

# Review: making prediction

$$\hat{\pi}_c = \frac{m_c}{m}$$

$$\hat{\theta}_{cd} = \frac{m_{cd}}{m_c}$$

Given one data point, it has 4 features (input), compare the probabilities:

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 0) &= p(y = 0) \cdot p(x_1, x_2, x_3, x_4 | y = 0) \\ &= p(y = 0) \cdot p(x_1 | y = 0) \cdot p(x_2 | y = 0) \cdot p(x_3 | y = 0) \cdot p(x_4 | y = 0) \end{aligned}$$

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 1) &= p(y = 1) \cdot p(x_1, x_2, x_3, x_4 | y = 1) \\ &= p(y = 1) \cdot p(x_1 | y = 1) \cdot p(x_2 | y = 1) \cdot p(x_3 | y = 1) \cdot p(x_4 | y = 1) \end{aligned}$$

# Bernoulli Naïve Bayes MLE: issue

*no data points of class 2:*

$$p(y = 2) = 0$$

*no data points in class 1 & 3 is 0 for x1:*

$$p(x_1 = 0|y = 3) = 0$$

$$p(x_1 = 0|y = 1) = 0$$

$y$	$x_1$	$x_2$
1	1	1
3	1	0
3	1	1
3	1	0
1	1	0
?	0	0

What if there are *no* examples of class  $c$  in the training set?

$$\hat{\pi}_c = 0 \quad \text{Model will never learn to guess class } c$$

What if all data points  $i$  in class  $c$  has  $x_d^{(i)} = 0$  in the training set?

$$\hat{\theta}_{cd} = 0$$

Model will assign 0 likelihood for test data with  $x_d = 1$  for class  $c$  (i.e.,  $p(x|y = c)$  ).

What does it imply on  $p(y = c|x)$  ? 0!

Training data needs to see every possible outcome for each feature

**Any ideas how we can fix this problem?**

We could add a small constant to prevent zero probabilities...

$$\hat{\pi}_c \propto m_c + \alpha$$

$$\hat{\theta}_{cj} \propto m_{cj} + \beta$$

$$\alpha, \beta > 0$$

**Pseudocounts  
add- $\alpha$  Smoothing  
Laplace smoothing**

....

**typical choice: set  $\alpha = \beta = 1$**

Another smoothing method:

$$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c) + 1}{\sum_{w \in V} (\text{count}(w, c) + 1)} = \frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|}$$

**Word count in category c**

**Vocabulary size  
in whole corpus**

# Naïve Bayes in Sentiment Classification

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	predictable with no fun

$$\hat{\pi}_c = \frac{m_c}{m} \quad P(-) = \frac{3}{5} \quad P(+) = \frac{2}{5}$$

# Naïve Bayes in Sentiment Classification

$$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c)}{\sum_{w \in V} \text{count}(w, c)}$$

smoothing  
↓

$$\frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|}$$

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	predictable with no fun

Vocabulary size

$$20 = 14 + 9 - 3$$

3: the, and, very (duplicate)

$$P(\text{"predictable"}|-) = \frac{1+1}{14+20}$$

$$P(\text{"predictable"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"no"}|-) = \frac{1+1}{14+20}$$

$$P(\text{"no"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"fun"}|-) = \frac{0+1}{14+20}$$

$$P(\text{"fun"}|+) = \frac{1+1}{9+20}$$



# Naïve Bayes in Sentiment Classification

	Cat	Documents
Training	-	just plain boring
	-	entirely predictable and lacks energy
	-	no surprises and very few laughs
	+	very powerful
	+	the most fun film of the summer
Test	?	predictable with no fun

Ignore unknown words

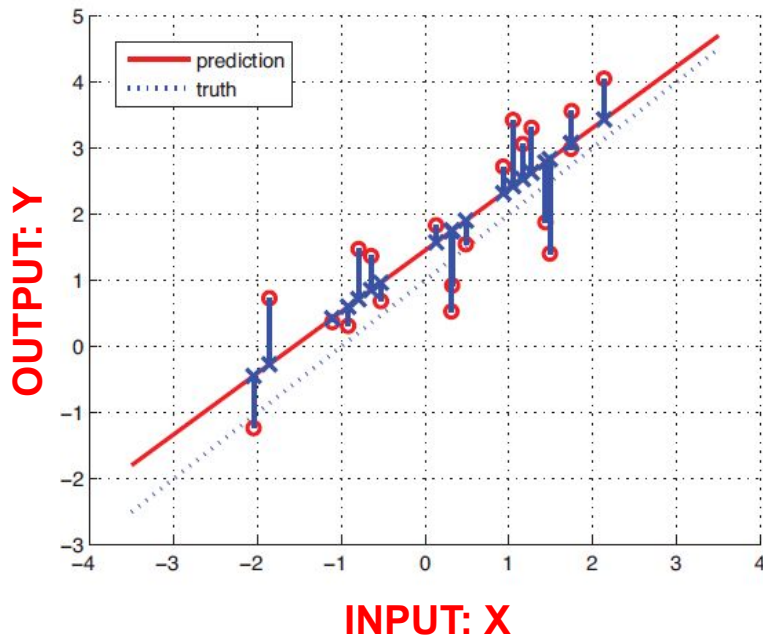


$$P(-)P(S|-) = \frac{3}{5} \times \frac{2 \times 2 \times 1}{34^3} = 6.1 \times 10^{-5}$$

$$P(+)P(S|+) = \frac{2}{5} \times \frac{1 \times 1 \times 2}{29^3} = 3.2 \times 10^{-5}$$

The model thus predicts the class *negative* for the test sentence.

# Linear Regression



**Regression** Learn a function that predicts outputs from inputs,

$$y = f(x)$$

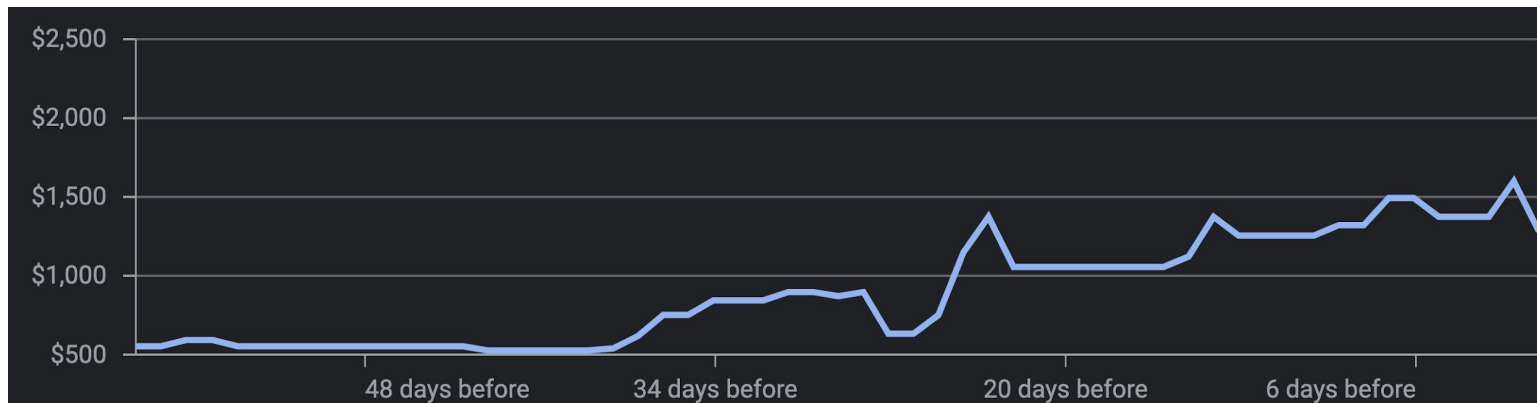
Outputs  $y$  are real-valued

**Linear Regression** As the name suggests, uses a *linear function*:

$$y = w^T x + b$$

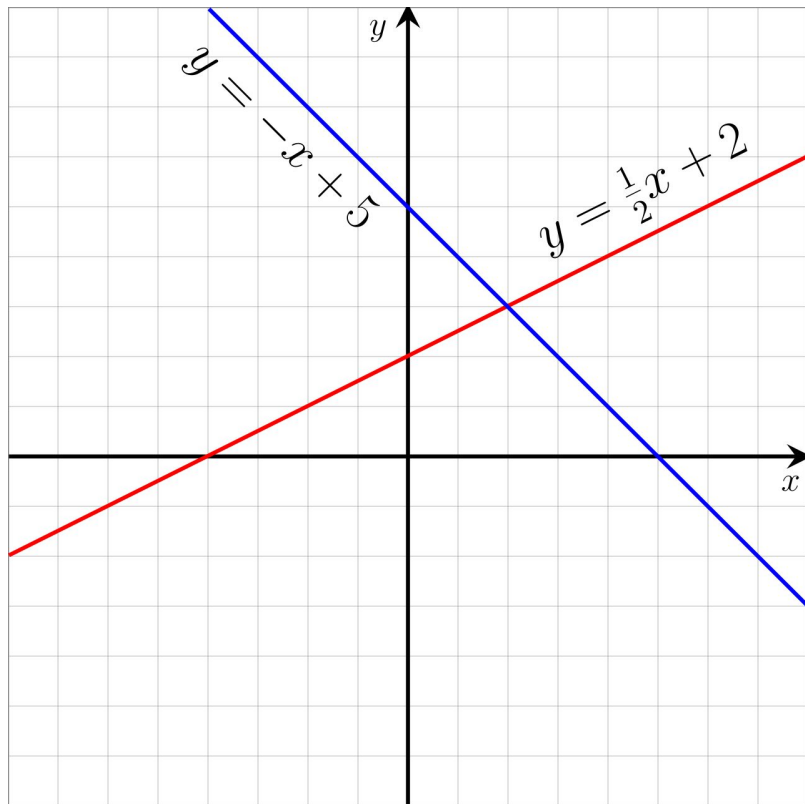
$$w^T x := \sum_{d=1}^D w_d x_d$$

## When is linear regression useful?



Price of an airline ticket

*Used anywhere a linear relationship is assumed between inputs / (real-valued) outputs*



Recall the equation for a line has a *slope* and an *intercept*,

$$y = w \cdot x + b$$

**Slope**      **Intercept**

- Intercept (b) indicates where line crosses y-axis
- Slope controls angle of line
- Positive slope (w) → Line goes up left-to-right
- Negative slope → Line goes down left-to-right

# Review: inner product

Two vectors:

$$\vec{x} = \langle 2, -3 \rangle \quad \mathbf{x} = \begin{bmatrix} 2 \\ -3 \end{bmatrix}$$

$$\vec{y} = \langle 5, 1 \rangle \quad \mathbf{y} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}$$

Multiply corresponding entries and add:

$$\vec{x} \cdot \vec{y} = \langle 2, -3 \rangle \cdot \langle 5, 1 \rangle = (2)(5) + (-3)(1) = 7$$

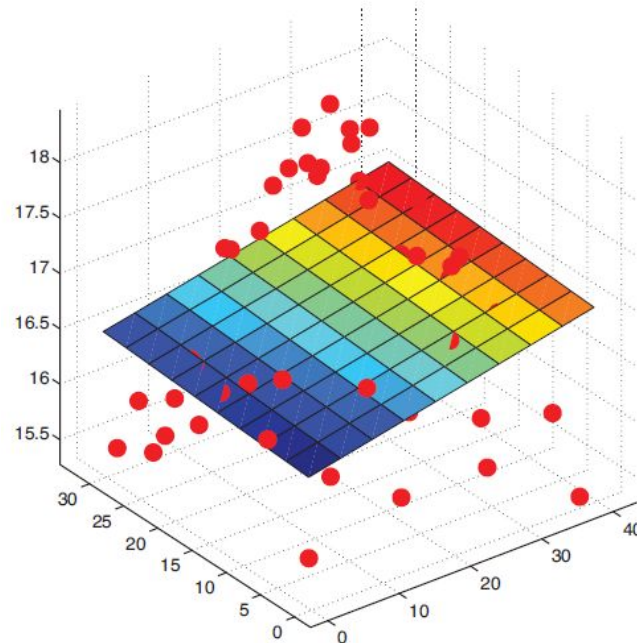
$$\mathbf{x}^T \mathbf{y} = \begin{bmatrix} 2 & -3 \end{bmatrix} \begin{bmatrix} 5 \\ 1 \end{bmatrix} = [7] \quad (\text{or just } 7) \quad (\text{so } \vec{x} \cdot \vec{y} \text{ becomes } \mathbf{x}^T \mathbf{y})$$

- **1d regression**: regression with 1d input:  
$$y = wx + b$$
- **D-dimensional regression**: input vector is  $x \in \mathbb{R}^D$ .

Recall the definition of an *inner product*:

$$w^T x = w_1 x_1 + w_2 x_2 + \dots + w_D x_D = \sum_{d=1}^D w_d x_d$$

The model is  $y = w^T x + b$



[ Image: Murphy, K. (2012) ]

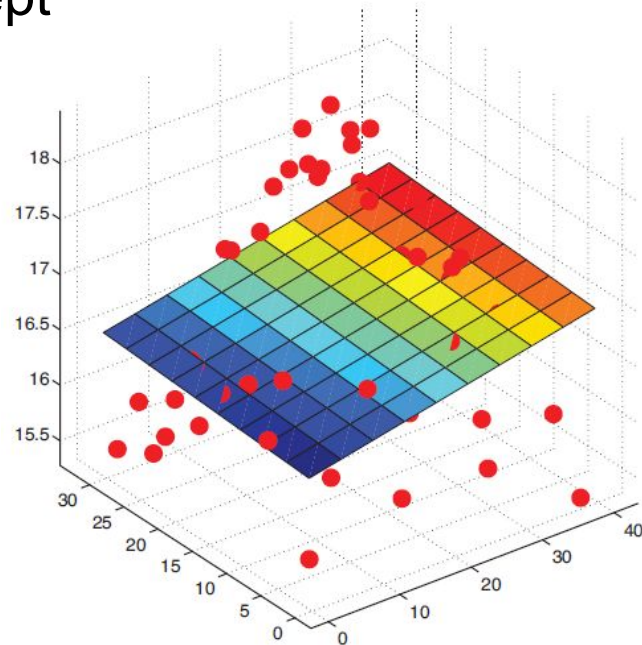
Often we simplify this by including the intercept into the weight vector,

$$\tilde{w} = \begin{pmatrix} w_1 \\ \vdots \\ w_D \\ b \end{pmatrix} \quad \tilde{x} = \begin{pmatrix} x_1 \\ \vdots \\ x_D \\ 1 \end{pmatrix} \quad y = \tilde{w}^T \tilde{x}$$

Since:

$$\begin{aligned} \tilde{w}^T \tilde{x} &= \sum_{d=1}^D w_d x_d + b \cdot 1 \\ &= w^T x + b \end{aligned}$$

from now on, we assume that  $w \in \mathbb{R}^D$  and  $x \in \mathbb{R}^D$  already has  $b$  and  $1$  in the last coordinate respectively.

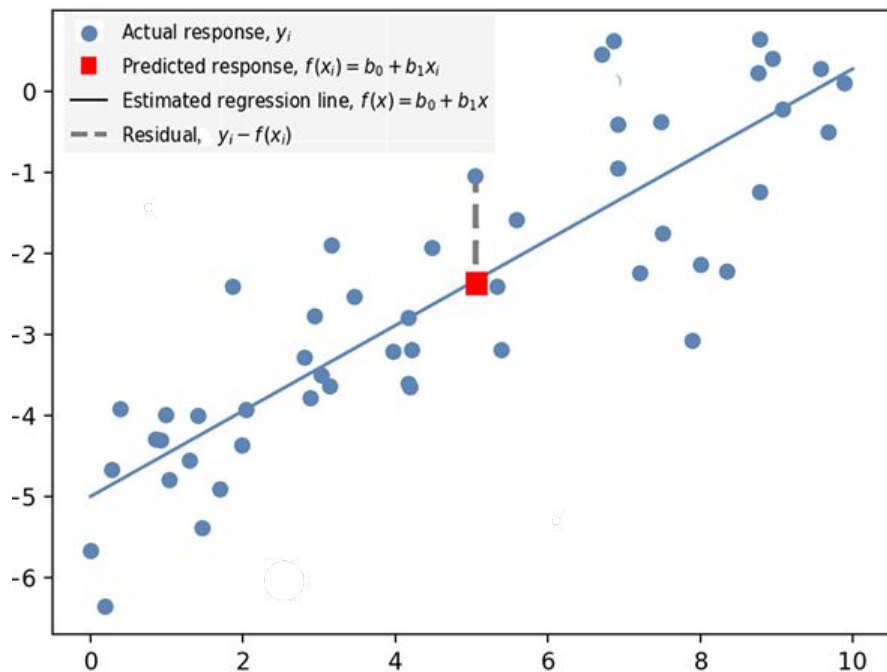




**There are several ways to think about fitting regression:**

- **Intuitive** Find a plane/line that is close to data
- **Functional** Find a line that minimizes the *least squares* loss
- **Estimation** Find maximum likelihood estimate of parameters

*They are all the same thing...*



**Intuition** Find a line that is as *close as possible* to every training data point

The distance from each point to the line is the **residual**

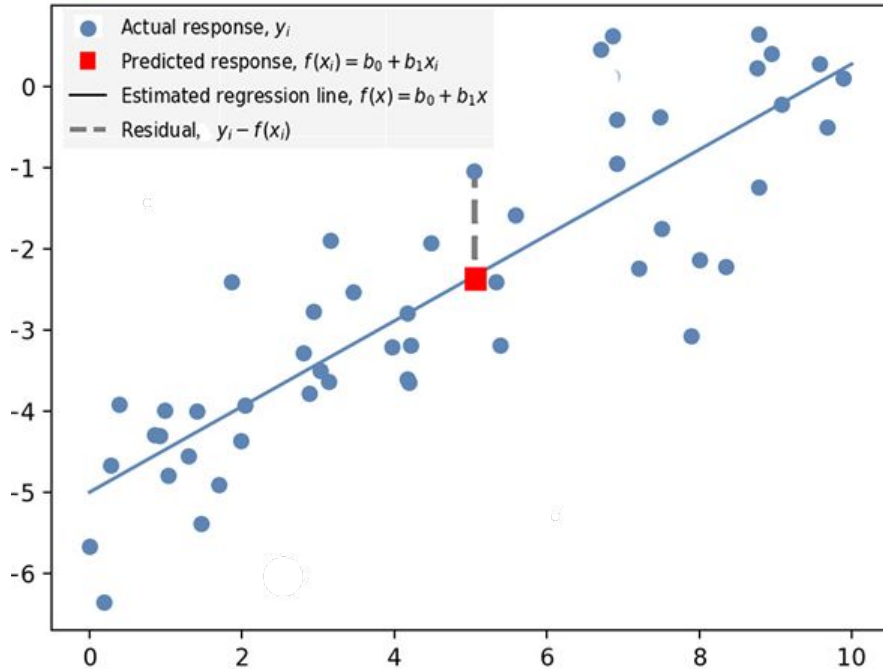
$$y - w^T x$$

Training Output

Prediction

*Let's find  $w$  that will minimize the residual!*

- Linear Regression
- **Least Squares Estimation**
- Regularized Least Squares
- Logistic Regression



**Functional** Find a line that minimizes the sum of squared residuals!

Given:  $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$

Compute:

$$w^* = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2$$

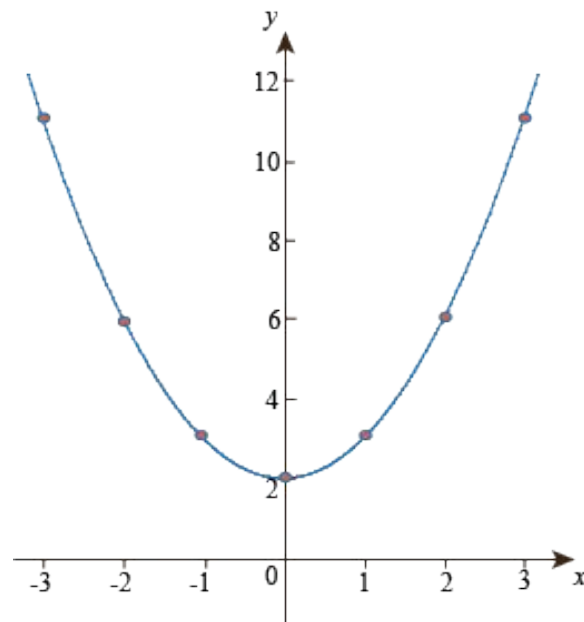
*Least squares regression*

$$\min_w \sum_{i=1}^N (y^{(i)} - w^T x^{(i)})^2$$

**This is just a quadratic function...**

- *Convex*, unique minimum
- Minimum given by zero-derivative
- Can find a closed-form solution

Let's see for scalar case with no bias,  
 $y = wx$



$$\frac{d}{dw} \sum_{i=1}^N (y^{(i)} - wx^{(i)})^2 =$$

**Derivative (+ chain rule)**

$$= \sum_{i=1}^N 2(y^{(i)} - wx^{(i)})(-x^{(i)}) = 0 \Rightarrow$$

**Distributive Property  
(and multiply -1 both sides)**

$$0 = \sum_{i=1}^N y^{(i)} x^{(i)} - w \sum_{j=1}^N (x^{(j)})^2$$

**Algebra**

$$w = \frac{\sum_i y^{(i)} x^{(i)}}{\sum_j (x^{(j)})^2}$$

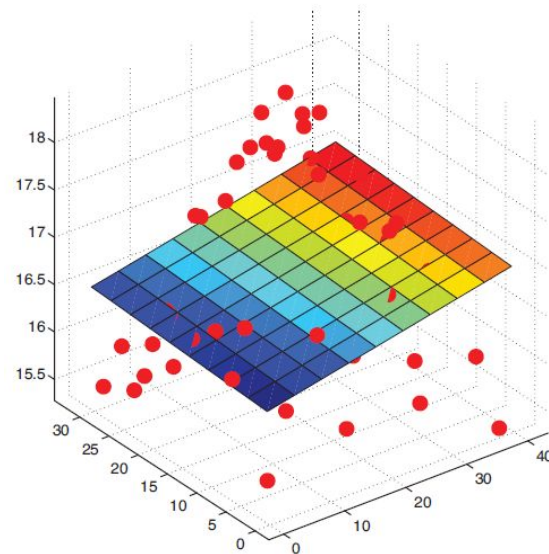
Things are a bit more complicated in higher dimensions and involve more linear algebra,

$$\mathbf{X} = \begin{pmatrix} x_1^{(1)} & \dots & x_D^{(1)} & 1 \\ x_1^{(2)} & \dots & x_D^{(2)} & 1 \\ \vdots & \vdots & \vdots & \vdots \\ x_1^{(m)} & \dots & x_D^{(m)} & 1 \end{pmatrix}$$

**Design Matrix**  
( each row is a data point)

$$\mathbf{y} = \begin{pmatrix} y^{(1)} \\ \vdots \\ y^{(N)} \end{pmatrix}$$

**Vector of labels**



Can write regression over *all training data* more compactly...

$$\mathbf{y} \approx \mathbf{X}\mathbf{w}$$

← **mx1 Vector**

$$= \begin{pmatrix} (x^{(1)})^\top \mathbf{w} \\ \vdots \\ (x^{(m)})^\top \mathbf{w} \end{pmatrix}$$

Least squares can also be written more compactly,

$$\|x\| := \sqrt{x \cdot x}.$$

$$\min_w \sum_{i=1}^N (y^{(i)} - w^T x^{(i)})^2 = \|\mathbf{y} - \mathbf{X}w\|^2$$

Some slightly more advanced linear algebra gives us a solution,

$$w = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \text{compare with the 1d version: } w = \frac{\sum_i y^{(i)} x^{(i)}}{\sum_j (x^{(j)})^2}$$

**Ordinary Least Squares (OLS)** solution

Derivation a bit advanced for this class, but enough to know

- it has a closed-form and why
- we can evaluate it
- generally know where it comes from.

