




Computer
Science

CSC380: Principles of Data Science

Linear Models 2

Xinchen Yu



Functional Find a line that minimizes the sum of squared residuals!

Given: $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$

Compute:

$$w^* = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2$$

Least squares regression

Least squares can also be written more compactly,


$$\|x\| := \sqrt{x \cdot x}.$$

$$\min_w \sum_{i=1}^N (y^{(i)} - w^T x^{(i)})^2 = \|\mathbf{y} - \mathbf{X}w\|^2$$

Some slightly more advanced linear algebra gives us a solution,

$$w = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Ordinary Least Squares (OLS) solution




There are several ways to think about fitting regression:

- **Intuitive** Find a plane/line that is close to data
- **Functional** Find a line that minimizes the *least squares* loss
- **Estimation** Find maximum likelihood estimate of parameters

They are all the same thing...

Probabilistic Assumptions



- Assume $x \sim \mathcal{D}_X$ from some distribution. We then assume that

$$y = w^T x + \epsilon \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2)$$

- Assume $x \sim \mathcal{D}_X$ from some distribution. We then assume that

$$y = w^T x + \epsilon \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2)$$

- Equivalently,

$$p(y|x; w) = \mathcal{N}(w^T x, \sigma^2)$$

Why? Adding a constant to a Normal RV is still a Normal RV,

$$z \sim \mathcal{N}(m, P) \qquad z + c \sim \mathcal{N}(m + c, P)$$

for our case, linear regression $z \leftarrow \epsilon$ and $c \leftarrow w^T x$

Given training data $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$, maximize the likelihood!

$$\hat{w} = \arg \max_w \log \prod_{i=1}^m p(x^{(i)}, y^{(i)}; w)$$

$$= \arg \max_w \log \prod_{i=1}^m p(x^{(i)}) p(y^{(i)} | x^{(i)}; w)$$

note $p(x^{(i)})$ does not depend on w !

$$= \arg \max_w \log \prod_{i=1}^m p(y^{(i)} | x^{(i)}; w)$$

subtracting a constant w.r.t. w does not affect the solution w !

$$= \arg \max_w \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}; w)$$

note model assumption! $p(y|x; w) = \mathcal{N}(w^T x, \sigma^2)$

Let's focus on 1d case.

Let $\mu = w^T x$ for now.

Gaussian (a.k.a. Normal) distribution with mean (location) μ and variance (squared scale) σ^2 parameters,



$$\mathcal{N}(y; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2}(y - \mu)^2 / \sigma^2\right)$$


The logarithm of the PDF is just a negative quadratic,

$$\log \mathcal{N}(y; \mu, \sigma^2) = \underbrace{-\frac{1}{2} \log 2\pi - \log \sigma}_{\text{Constant w.r.t. mean}} - \underbrace{\frac{1}{2\sigma^2} (y - \mu)^2}_{\text{Quadratic Function of mean}}$$

Constant w.r.t. mean

Quadratic Function of mean





Substitute linear regression prediction into MLE solution and we have,

$$\arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2$$

So for Linear Regression,
MLE = Least Squares Estimation

1. The linear regression model (assumption),

$$y = w^T x + \epsilon \quad \text{where} \quad \epsilon \sim \mathcal{N}(0, \sigma^2 I)$$

2. For N iid training data fit using least squares,

$$w^{\text{OLS}} = \arg \min_w \sum_{i=1}^N (y^{(i)} - w^T x^{(i)})^2$$

3. Equivalent to maximum likelihood solution

$$w^{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Least squares solution requires inversion of the term,

$$(\mathbf{X}^T \mathbf{X})^{-1}$$

What is the issue?

May be non-invertible!

Invertible matrix

Invertible matrix: a matrix A of dimension $n \times n$ is called invertible if and only if there exists another matrix B of the same dimension, such that $AB = BA = I$, where I is the identity matrix of the same order.

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 5 \end{bmatrix} \quad AB = \begin{bmatrix} 1 & 2 \\ 2 & 5 \end{bmatrix} \begin{bmatrix} 5 & -2 \\ -2 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$
$$B = \begin{bmatrix} 5 & -2 \\ -2 & 1 \end{bmatrix} \quad BA = \begin{bmatrix} 5 & -2 \\ -2 & 1 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 2 & 5 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$w^{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Use *Moore-Penrose pseudoinverse* ('dagger' notation)

$$w^{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{\dagger} \mathbf{X}^T \mathbf{y}$$

- Generalization of the standard matrix inverse for non-invertible matrices.
- Directly computable in most libraries
- In Numpy it is: `linalg.pinv`



Load your libraries,

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score
```

For Evaluation



Load data,

```
# Load the diabetes dataset
diabetes_X, diabetes_y = datasets.load_diabetes(return_X_y=True)

# Use only one feature
diabetes_X = diabetes_X[:, np.newaxis, 2]
```

^: same as `diabetes_X[:,2]`

Samples total	442
Dimensionality	10
Features	real, $-0.2 < x < 0.2$
Targets	integer 25 - 346

Train / Test Split:

```
diabetes_X_train = diabetes_X[:-20]
diabetes_X_test = diabetes_X[-20:]
```

```
diabetes_y_train = diabetes_y[:-20]
diabetes_y_test = diabetes_y[-20:]
```

Train (fit) and predict,

```
# Create linear regression object
regr = linear_model.LinearRegression()

# Train the model using the training sets
regr.fit(diabetes_X_train, diabetes_y_train)

# Make predictions using the testing set
diabetes_y_pred = regr.predict(diabetes_X_test)
```

Plot regression line with the test set,

```
# Plot outputs
plt.scatter(diabetes_X_test, diabetes_y_test, color="black")
plt.plot(diabetes_X_test, diabetes_y_pred, color="blue", linewidth=3)

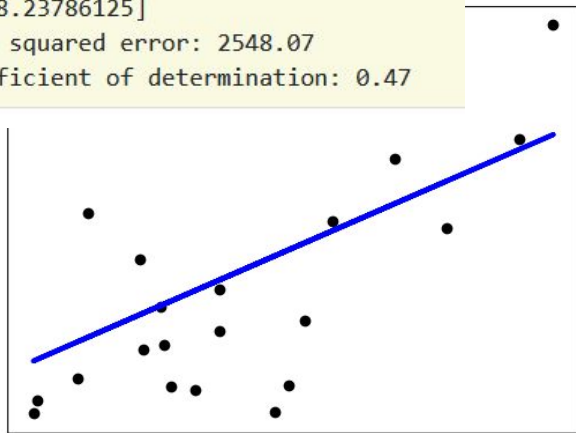
plt.xticks(())
plt.yticks(())

plt.show()
```

`regr.coef_` : coefficient (array)
`regr.intercept_` : intercept (float)



Coefficients:
[938.23786125]
Mean squared error: 2548.07
Coefficient of determination: 0.47



Regularized Least Squares

Recall: OLS solution

$$w^{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Use *Moore-Penrose pseudoinverse* ('dagger' notation)

$$w^{\text{OLS}} = (\mathbf{X}^T \mathbf{X})^{\dagger} \mathbf{X}^T \mathbf{y}$$

Or, use L2 Regularized Least Squares (RLS)

$$w^{\text{L2}} = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

Q: why is this called regularized least squares?

$$w^{\text{L2}} = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

Turns out, w^{L2} is the solution of

$$w^{\text{L2}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 + \lambda \|w\|^2 \quad \text{recall: } \|w\| = \sqrt{\sum_{d=1}^D w_d^2}$$



λ : Regularization Strength

$\|w\|^2$: Regularization Penalty

Prefers smaller magnitudes for w !

λ very small: almost OLS


λ very large: $w \approx 0$ and high trainset error

Okay, we have a training data. Why not learn the most complex function that can work flawlessly for the training data and be done with it? (i.e., classifies every data point correctly)

Extreme example: Let's memorize the data. To predict an unseen data, just follow the label of the closest memorized data.

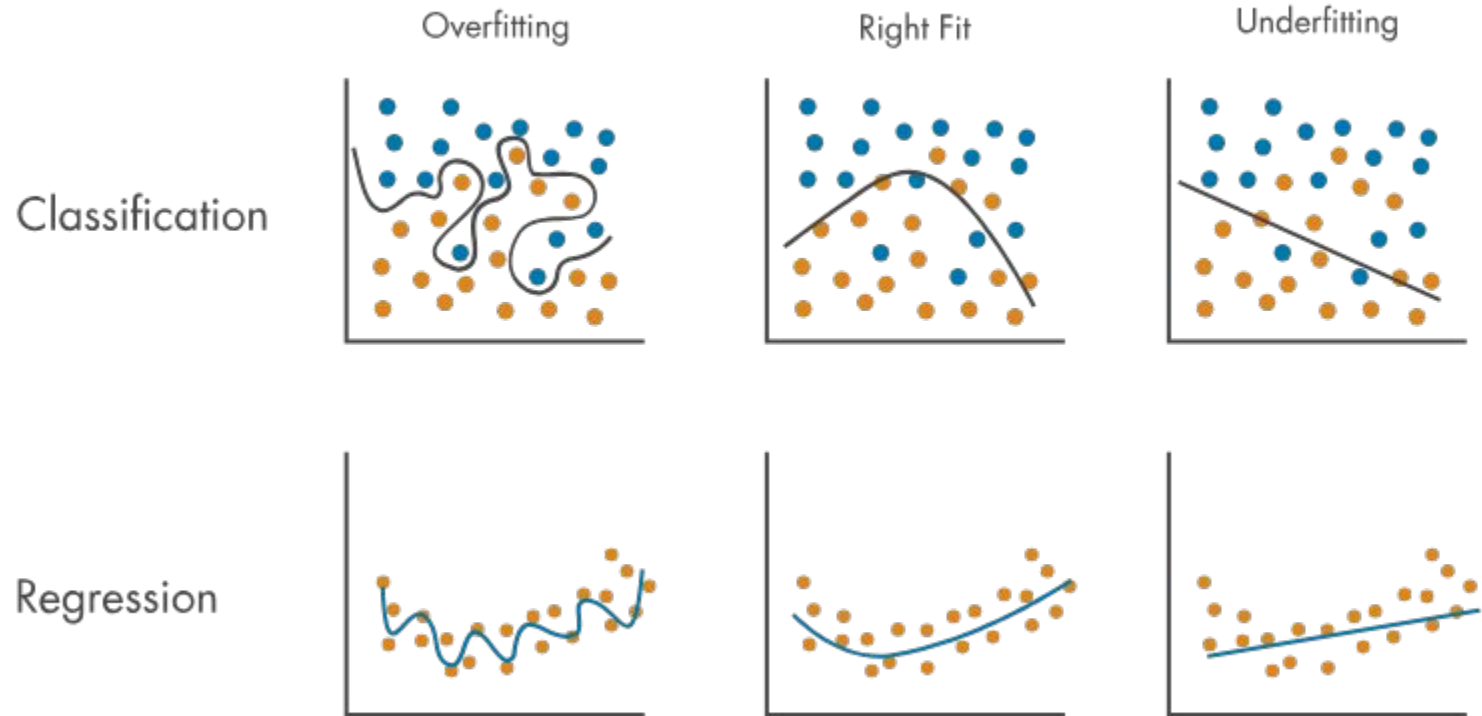
Doesn't generalize to unseen data – called *overfitting* the training data.

Solution: Fit the train set but don't "over-do" it. This is called **regularization**.




green: almost memorization
black: true decision boundary

Overfitting vs Underfitting

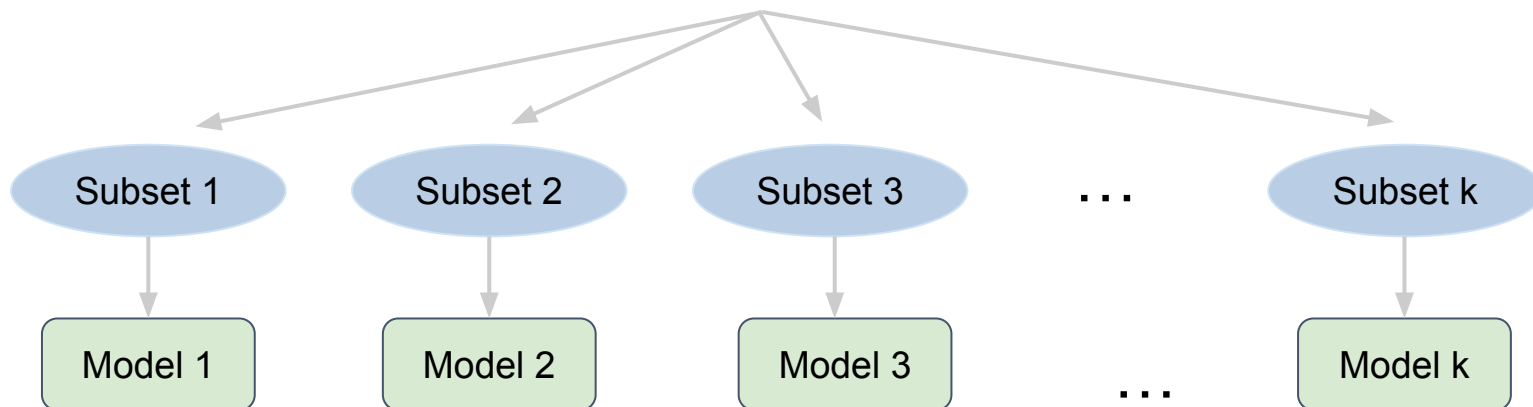


Bias-Variance Tradeoff




Bias-Variance Tradeoff: an example

Student ID	GPA	Working experiences	...	Demographics
0
...



Task: given a new student, predict if accepted or rejected.

Option 1: overly simple model



Models are similar but wrong on average:

- Bias high
- Variance low

\overline{GPA}_i is similar for $i = 1, 2, \dots, k$

Model 1


Model 2

Model 3

...

Model k

Option 2: overly complex model



Models are different but right on average:

- Bias low
- Variance high

small changes in training set → big changes in the model

Model 1

Model 2

Model 3

...

Model k

Regularization

- 1d case

- Suppose that $y = wx + \epsilon$, and the true model is $w = 0$ ($y = \epsilon$)
- However, OLS is highly probable to 'exaggerate' the effect of x to decrease train set error: (overfitting)

$$w = \frac{\sum_i y^{(i)} x^{(i)}}{\sum_j (x^{(j)})^2}$$

- On the other hand, RLS will try to balance the train set error and the penalty caused by the large norm

$$w^{RLS} = \frac{\sum_i y^{(i)} x^{(i)}}{\sum_j (x^{(j)})^2 + \lambda}$$
$$|w^{RLS}| < |w^{OLS}|$$

$$w^{\text{RLS}} = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

Turns out, w^{RLS} is the solution of

$$w^{\text{L2}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 + \lambda \|w\|^2 \quad \text{recall: } \|w\| = \sqrt{\sum_{d=1}^D w_d^2}$$



λ : Regularization Strength

$\|w\|^2$: Regularization Penalty

In short, the benefits of L2-RLS

- No need to worry about the estimator being undefined (due to matrix inversion)
- Avoid overfitting (if λ is chosen well)!

`sklearn.linear_model.Ridge`

```
class sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, normalize='deprecated', copy_X=True, max_iter=None, tol=0.001, solver='auto', positive=False, random_state=None) ⓘ
```

[\[source\]](#)

Minimizes the objective function:

$$||y - Xw||^2_2 + \alpha * ||w||^2_2$$

Alpha is what we have been calling λ

alpha : {float, ndarray of shape (n_targets,)}, default=1.0


Regularization strength; must be a positive float. Regularization improves the conditioning of the problem and reduces the variance of the estimates. Larger values specify stronger regularization. Alpha corresponds to $1 / (2C)$ in other linear models such as `LogisticRegression` or `LinearSVC`. If an array is passed, penalties are assumed to be specific to the targets. Hence they must correspond in number.

Define and fit OLS and L2 regression,

```
ols=linear_model.LinearRegression()  
ols.fit(X_train, y_train)  
ridge=linear_model.Ridge(alpha=0.1)  
ridge.fit(X_train, y_train)
```

Plot results,

```
fig, ax = plt.subplots()  
ax.scatter(X_train, y_train, s=50, c="black", marker="o")  
ax.plot(X_test, ols.predict(X_test), color="red", label="OLS")  
ax.plot(X_test, ridge.predict(X_test), color="blue", label="L2")  
  
plt.legend()  
plt.show()
```



quiz candidate

Q: why L2 has a lower slope?

L2 (Ridge) reduces impact of any single data point

- Feature weights are “shrunk” towards zero – statisticians often call this a “shrinkage” method
- Common practice: Do **not** penalize bias (y-intercept, w_D) parameter,


$$\min_w \sum_i (y^{(i)} - w^T x^{(i)})^2 + \frac{\lambda}{2} \sum_{d=1}^{D-1} w_d^2$$

Recall: we enforced $x_D^{(i)} = 1$ so that w_D encodes the intercept

- Penalizing intercept will make solution depend on origin for Y.
i.e., add a constant c to $y^{(i)}$'s \Rightarrow the solutions changes!

We need to tune regularization strength to avoid over/under fitting...

$$w^{\text{L2}} = \arg \min_w \sum_{i=1}^m (y^{(i)} - w^T x^{(i)})^2 + \lambda \|w\|^2$$



Recall bias/variance tradeoff

High regularization *reduces* model complexity: *increases* bias / *decreases* variance

Q: How should we properly tune λ ?

cross validation!