



Computer
Science

CSC380: Principles of Data Science

Predictive Modeling and Classification 3

Xinchen Yu

9. Given a distribution D with unknown mean μ and variance σ^2 , and a set of n iid samples X_1, \dots, X_n drawn from it. Define $\tilde{\mu}_n = \frac{1}{n-1} \sum_{i=1}^n X_i$ as an estimator of μ .

(a) (4 points) Is $\tilde{\mu}_n$ an unbiased estimator of μ ? Justify your answer.

(b) (6 points) Let $n = 4$. What is the bias, variance, and Mean Square Error (MSE) of $\tilde{\mu}_4$, respectively? Note: For variance, you can compute $Var[\tilde{\mu}_4]$, in other words, $Var[\frac{X_1+X_2+X_3+X_4}{3}]$.

(You can have μ, σ^2 or numbers in the results).

Lecture statistics 3, page 7

$$\begin{aligned}\tilde{\mu}_n &= \frac{1}{n-1} \sum_{i=1}^n X_i & E[\tilde{\mu}_n] &= E\left[\frac{1}{n-1} \sum_{i=1}^n X_i\right] \\ & & &= \frac{1}{n-1} E\left[\sum_{i=1}^n X_i\right] \\ & & &= \frac{1}{n-1} \sum_{i=1}^n E[X_i] \\ & & &= \frac{1}{n-1} \sum_{i=1}^n \mu = \frac{n\mu}{n-1}\end{aligned}$$

$\tilde{\mu}_n$ is not an unbiased estimator of μ .

9. Given a distribution D with unknown mean μ and variance σ^2 , and a set of n iid samples X_1, \dots, X_n drawn from it. Define $\tilde{\mu}_n = \frac{1}{n-1} \sum_{i=1}^n X_i$ as an estimator of μ .

(a) (4 points) Is $\tilde{\mu}_n$ an unbiased estimator of μ ? Justify your answer.

(b) (6 points) Let $n = 4$. What is the bias, variance, and Mean Square Error (MSE) of $\tilde{\mu}_4$, respectively? Note: For variance, you can compute $\text{Var}[\tilde{\mu}_4]$, in other words, $\text{Var}[\frac{X_1+X_2+X_3+X_4}{3}]$.

(You can have μ, σ^2 or numbers in the results).

$$\tilde{\mu}_4 = \frac{1}{3}(X_1 + X_2 + X_3 + X_4) \quad \text{Var}[\tilde{\mu}_4] = \text{Var}\left[\frac{1}{3}(X_1 + X_2 + X_3 + X_4)\right]$$

$$\text{Bias}(\tilde{\mu}_4) = E[\tilde{\mu}_4] - \mu \quad = \frac{1}{9} \text{Var}[X_1 + X_2 + X_3 + X_4]$$

$$= \frac{4\mu}{3} - \mu$$

$$= \frac{\mu}{3}$$

Since the X_i are iid:

$$= \frac{1}{9}(\text{Var}[X_1] + \text{Var}[X_2] + \text{Var}[X_3] + \text{Var}[X_4])$$

$$= \frac{1}{9}(4\sigma^2)$$

$$\text{MSE}(\tilde{\mu}_4) = \text{Var}[\tilde{\mu}_4] + \text{Bias}(\tilde{\mu}_4)^2 = \frac{4\sigma^2 + \mu^2}{9}$$

Model Evaluation

Suppose our classifier distinguishes between cats and non-cats.
We can make the following table called **confusion matrix**:

Actual class \ Predicted class	Cat	Non-cat
	Cat	Non-cat
Cat	6 true positives	2 false negatives
Non-cat	1 false positive	3 true negatives

It tells us if classifier is biased towards certain mistakes (False Positives, False Neg.)

Good for investigating opportunities to improve the classifier.

		PREDICTED	
		POSITIVE	NEGATIVE
ACTUAL	POSITIVE	TRUE POSITIVES	FALSE NEGATIVES
	NEGATIVE	FALSE POSITIVES	TRUE NEGATIVES

Precision: dividing the true positives by anything that was predicted as a positive.

$$\frac{\text{TRUE POSITIVES}}{\text{TRUE POSITIVES} + \text{FALSE POSITIVES}}$$

Evaluating Classifiers - Recall

7

		PREDICTED	
		POSITIVE	NEGATIVE
ACTUAL	POSITIVE	TRUE POSITIVES	FALSE NEGATIVES
	NEGATIVE	FALSE POSITIVES	TRUE NEGATIVES

Recall (or True Positive Rate): dividing the true positives by anything that should have been predicted as positive.

$$\frac{\text{TRUE POSITIVES}}{\text{TRUE POSITIVES} + \text{FALSE NEGATIVES}}$$

F1 score symmetrically represents both precision and recall in one metric.




$$F_1 = \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}} = \frac{\text{tp}}{\text{tp} + \frac{1}{2}(\text{fp} + \text{fn})}$$

- This is the *harmonic mean* of precision and recall
 - `harmonic_mean(x,y)`

$$\frac{1}{\frac{1}{2}(\frac{1}{x} + \frac{1}{y})}$$

- Gives equal importance to precision and recall – F1 may not be best when you care about one more than the other (e.g., in medical tests we care about recall)

Evaluation functions live in `metrics`

<code>metrics.confusion_matrix(y_true, y_pred, *)</code>	Compute confusion matrix to evaluate the accuracy of a classification.	
<code>metrics.dcg_score(y_true, y_score, *[k, ...])</code>	Compute Discounted Cumulative Gain.	
<code>metrics.det_curve(y_true, y_score[, ...])</code>	Compute error rates for different probability thresholds.	
<code>metrics.f1_score(y_true, y_pred, *[, ...])</code>	Compute the F1 score, also known as balanced F-score or F-measure.	
<code>metrics.fbeta_score(y_true, y_pred, *, beta)</code>	Compute the F-beta score.	
<code>metrics.hamming_loss(y_true, y_pred, *[, ...])</code>	Compute the average Hamming loss.	
<code>metrics.hinge_loss(y_true, pred_decision, *)</code>	Average hinge loss (non-regularized).	
<code>metrics.jaccard_score(y_true, y_pred, *[, ...])</code>	Jaccard similarity coefficient score.	
<code>metrics.log_loss(y_true, y_pred, *[, eps, ...])</code>	Log loss, aka logistic loss or cross-entropy loss.	
<code>metrics.matthews_corrcoef(y_true, y_pred, *)</code>	Compute the Matthews correlation coefficient (MCC).	
<code>metrics.multilabel_confusion_matrix(y_true, ...)</code>	Compute a confusion matrix for each class or sample.	
<code>metrics.ndcg_score(y_true, y_score, *[k, ...])</code>	Compute Normalized Discounted Cumulative Gain.	
<code>metrics.precision_recall_curve(y_true, ...)</code>	Compute precision-recall pairs for different probability thresholds.	
<code>metrics.precision_recall_fscore_support(...)</code>	Compute precision, recall, F-measure and support for each class.	
<code>metrics.precision_score(y_true, y_pred, *[, ...])</code>	Compute the precision.	
<code>metrics.recall_score(y_true, y_pred, *[, ...])</code>	Compute the recall.	

Naïve Bayes

Heads Up This section will return to some math as we go in depth. But, much of it is that you already know with a new application (Naïve Bayes Classification) – **ask questions if you are lost**

- Introduction to Naïve Bayes Classifier
- Maximum Likelihood Estimation

Math Prep

- N RVs conditionally independent, given Z, if and only if:

$$p(X_1, \dots, X_N \mid Z) = \prod_{i=1}^N p(X_i \mid Z) \quad \text{Probability 3}$$

- Bayes' rule

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B|A)P(A)}{P(B)} \quad \text{Probability 2}$$

- Maximum likelihood estimation Statistics 2
- Bernoulli and Gaussian distribution Probability 3, 4
Statistics 2

Intuition

Training Data:

Person	height (feet)
male	6
male	5.92 (5'11")
male	5.58 (5'7")
male	5.92 (5'11")
female	5
female	5.5 (5'6")
female	5.42 (5'5")
female	5.75 (5'9")

Y **X**

Testing Data:

Person	height (feet)
?	6

Task: observe feature x_n , predict label $y_n \in (0, 1)$

Choose the class with higher probability:

$$P(Y_n = 1|X_n = x_n) \text{ vs } P(Y_n = 0|X_n = x_n)$$

$$P(Y_n = 1|X_n = x_n) = \frac{P(Y_n = 1, X_n = x_n)}{P(X_n = x_n)}$$

$$P(Y_n = 0|X_n = x_n) = \frac{P(Y_n = 0, X_n = x_n)}{P(X_n = x_n)}$$

How to learn a joint probability model for $P(Y, X)$?

Training Data:

Person	height (feet)	weight (lbs)	foot size(inches)
male	6	180	12
male	5.92 (5'11")	190	11
male	5.58 (5'7")	170	12
male	5.92 (5'11")	165	10
female	5	100	6
female	5.5 (5'6")	150	8
female	5.42 (5'5")	130	7
female	5.75 (5'9")	150	9



Task: Observe features x_1, \dots, x_D and predict class label $y \in \{1, \dots, C\}$

Model: Assume that the feature x and its label y follows certain type of distribution \mathcal{D} with parameter θ .

$$(x, y) \stackrel{\text{i.i.d.}}{\sim} \mathcal{D}_\theta$$

Training Algorithm: Estimate θ e.g., MLE $\hat{\theta}$

To classify: Compute

$$\hat{y} = \arg \max_{c \in \{1, \dots, C\}} p(y = c \mid x; \hat{\theta})$$

what comes after semicolon is the parameter of the distribution

Training Data:

Person	height (feet)	weight (lbs)	foot size(inches)
male	6	180	12
male	5.92 (5'11")	190	11
male	5.58 (5'7")	170	12
male	5.92 (5'11")	165	10
female	5	100	6
female	5.5 (5'6")	150	8
female	5.42 (5'5")	130	7
female	5.75 (5'9")	150	9



Features

Task: Observe features x_1, \dots, x_D and predict class label $y \in \{1, \dots, C\}$

Naïve Bayes Model: Treat features as *conditionally independent* given class label,

$$p(x, y) = p(y)p(x|y) = p(y) \prod_{d=1}^D p(x_d | y)$$

build individual models for these

To classify a given instance x : Bayes rule!

$$p(y = c | x) = \frac{p(y = c)p(x | y = c)}{p(x)}$$

Key concept in Naïve Bayes

$$p(x, y) = p(y) p(x|y) = p(y) \prod_{d=1}^D p(x_d | y)$$

Class prior distribution

Class conditional distribution

Given one data point, it has 4 features (input), and the label is 0 (output)

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 0) &= p(y = 0) \cdot p(x_1, x_2, x_3, x_4 | y = 0) \\ &= p(y = 0) \cdot p(x_1 | y = 0) \cdot p(x_2 | y = 0) \cdot p(x_3 | y = 0) \cdot p(x_4 | y = 0) \end{aligned}$$

$$p(x, y) = p(y) \underset{\text{Class prior distribution}}{p(x|y)} = p(y) \prod_{d=1}^D \underset{\text{Class conditional distribution}}{p(x_d | y)}$$

Class prior distribution

Class conditional distribution

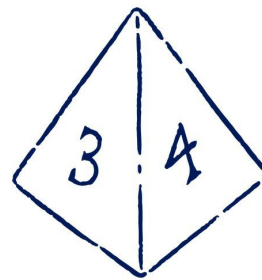
For the **class prior distribution**, take categorical distribution.

$$y \sim \text{Categorical}(\pi), \quad \pi \in \mathbb{R}^C, \pi_c \geq 0, \sum_c \pi_c = 1$$

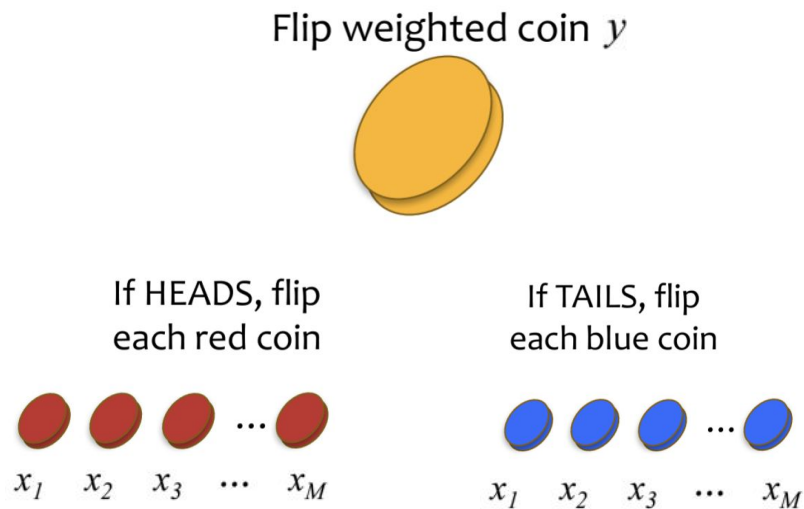
$$\Rightarrow p(y = c) = \pi_c$$

Example: biased 4-sided die Y, given:

- $P(Y=1) = 0.2, \quad P(Y=2) = 0.3, \quad P(Y=3) = 0.1, \quad P(Y=4) = 0.4$




Class prior distribution



y	x_1	x_2	x_3	...	x_M
0	1	0	1	...	1
1	0	1	0	...	1
1	1	1	1	...	1
0	0	0	1	...	1
0	1	0	1	...	0
1	1	0	1	...	0

$p(y)$

Class conditional distribution

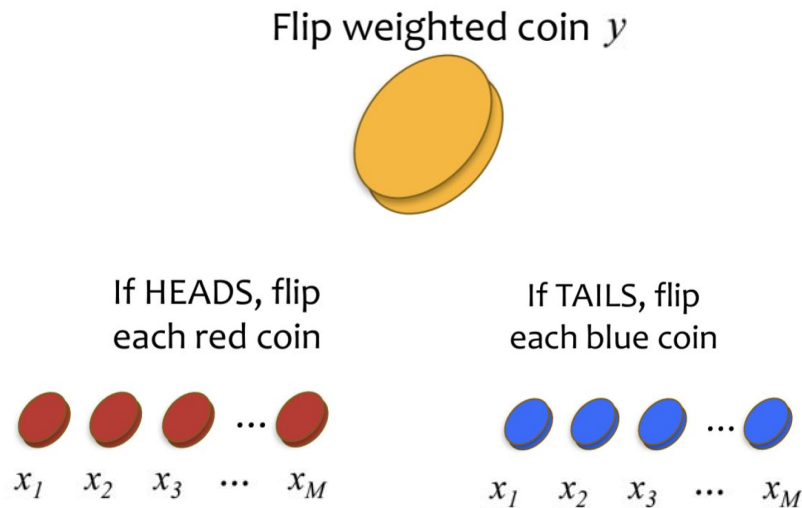


Each red / blue coin biases can be different.

y	x_1	x_2	x_3	...	x_M
0	1	0	1	...	1
1	0	1	0	...	1
1	1	1	1	...	1
0	0	0	1	...	1
0	1	0	1	...	0
1	1	0	1	...	0

$p(x_1|y=1)$

Class conditional distribution

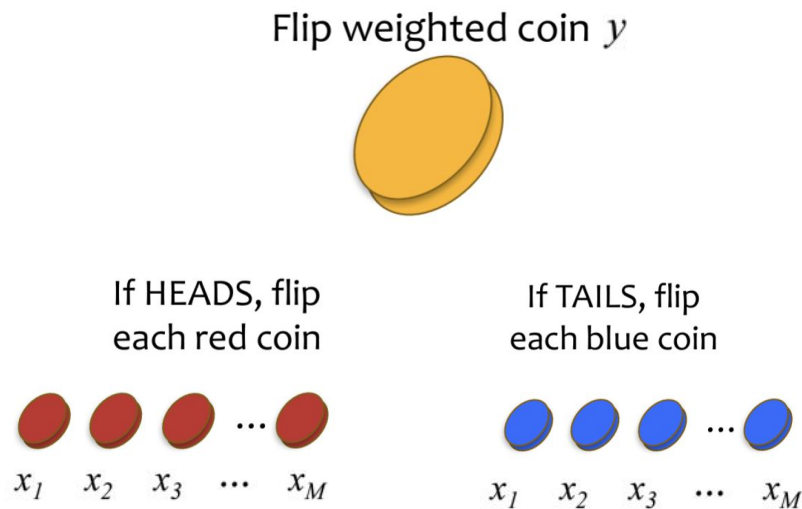


Each red / blue coin biases can be different.

y	x_1	x_2	x_3	...	x_M
0	1	0	1	...	1
1	0	1	0	...	1
1	1	1	1	...	1
0	0	0	1	...	1
0	1	0	1	...	0
1	1	0	1	...	0

$p(x_2|y = 1)$

Class conditional distribution



Each red / blue coin biases can be different.

y	x_1	x_2	x_3	...	x_M
0	1	0	1	...	1
1	0	1	0	...	1
1	1	1	1	...	1
0	0	0	1	...	1
0	1	0	1	...	0
1	1	0	1	...	0

\uparrow
 $p(x_1|y=0)$

Simplifying Assumption: “*Class conditional*” distribution assumes features are conditionally independent given class

$$p(x | y) = \prod_{d=1}^D p(x_d | y)$$

- “Naïve” as in general features are likely to be dependent.
- Every feature can have a different class-conditional distribution



$P(x_1|y)$ can be different from $P(x_2|y)$

Features are typically not independent!

Example 1 If a recent news article contains word “Donald” it is much more likely to contain the word “Trump”.

Example 2 If flower petal width is very large then petal length is also likely to be high.



Simplifying Assumption: “*Class conditional*” distribution assumes features are conditionally independent given class

$$p(x | y) = \prod_{d=1}^D p(x_d | y)$$

- “Naïve” as in general features are likely to be dependent.
- Every feature can have a different class-conditional distribution

Doesn't capture correlation among features. But why would it be a good idea?

- Easy computation: For C classes and D features only $O(CD)$ parameters
- Prevents overfitting
- Simplicity

For real-valued features we can use Normal distribution:

$$p(x | y = c) = \prod_{d=1}^D \mathcal{N}(x_d | \mu_{cd}, \sigma_{cd}^2)$$

quiz candidate

Q: how many parameters?

2cd

Parameters of featured for class c

For binary features $x_d \in \{0,1\}$ can use Bernoulli distributions:

$$p(x | y = c) = \prod_{d=1}^D \text{Bernoulli}(x_d | \theta_{cd})$$

quiz candidate

Q: how many parameters?

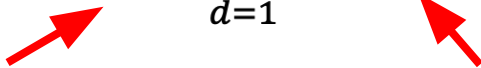
cd

“Coin bias” for d^{th} feature
and class c

- K-valued discrete features: use Categorical.
- Can mix-and-match, e.g. some discrete, some continuous features

$$p(x | y = c) = \prod_{d=1}^{D'} \text{Bernoulli}(x_d | \theta_{cd}) \prod_{d=D'+1}^D \mathcal{N}(x_d | \mu_{cd}, \sigma_{cd}^2)$$

Fitting the model requires learning all parameters...

$$p(x, y = c) = p(y = c; \pi) \prod_{d=1}^D p(x_d | \theta_{cd})$$


Class Prior Parameters

Likelihood Parameters

Given training data $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ maximize the likelihood function,

$$\theta^{\text{MLE}} = \arg \max_{\pi, \theta} \log p(\mathcal{D}; \pi, \theta)$$

$$\theta^{\text{MLE}} = \arg \max_{\pi, \theta} \log p(\mathcal{D}; \pi, \theta) \quad (\mathcal{D} := \{(x^{(i)}, y^{(i)})\}_{i=1}^m)$$

Since data are iid

$$= \arg \max_{\pi, \theta} \log \prod_{i=1}^m p(x^{(i)}, y^{(i)}; \pi, \theta)$$

$\log(ab) = \log a + \log b$

$$= \arg \max_{\pi, \theta} \sum_{i=1}^m \log p(x^{(i)}, y^{(i)}; \pi, \theta)$$

Conditional probability +
Naïve Bayes assumption

$$= \arg \max_{\pi, \theta} \sum_{i=1}^m \log p(y^{(i)}; \pi) + \sum_{i=1}^m \sum_{d=1}^D \log p(x_d^{(i)} | y^{(i)}; \theta_{y^{(i)}d})$$

θ_{cd} : parameter for feature d for class c

Find zero-gradient if concave, or gradient-based optimization otherwise

Analogy:

y	x_1	x_2	...	x_D		
5	0	1	1	0	1	0
2	1	0	0	0	0	0
3	1	1	1	1	0	0
...	...					
4	0	0	1	1	0	1

Let each feature follow a Bernoulli distribution then the model is...

$$y \sim \text{Categorical}(\pi) \quad x_d \mid y = c \sim \text{Bernoulli}(\theta_{cd})$$

The Naïve Bayes joint distribution is then:

$$\begin{aligned} p(\mathcal{D}; \pi, \theta) &= \prod_{i=1}^m \left(p(y^{(i)}; \pi) \prod_d p(x_d^{(i)}; \theta_{y^{(i)}d}) \right) \\ &= \prod_{i=1}^m \left(\prod_c \left(\pi_c^{\mathbf{I}\{y_i=c\}} \prod_j p(x_{ij} | \theta_{jc})^{\mathbf{I}\{y_i=c\}} \right) \right) \end{aligned}$$

Write down log-likelihood and optimize...

j : feature, c : label, i : data

$$y \sim \text{Categorical}(\pi_c) : p(y = c) = \pi_c$$

$$p(y = 1) = \pi_1$$

$$p(y = 2) = \pi_2$$

$$p(y = 3) = \pi_3 = 1 - \pi_1 - \pi_2$$

Q: how many parameters?

$c-1+cj$

$$x|y \sim \text{Bernoulli}(\theta_{jc}) : p(x|y) = \theta_{jc}^x (1 - \theta_{jc})^{1-x}$$

$$x_{j=1}|y = 1 \sim \text{Bernoulli}(\theta_{j=1,c=1}) \quad x_{j=2}|y = 1 \sim \text{Bernoulli}(\theta_{j=2,c=1})$$

$$x_{j=1}|y = 2 \sim \text{Bernoulli}(\theta_{j=1,c=2}) \quad x_{j=2}|y = 2 \sim \text{Bernoulli}(\theta_{j=2,c=2})$$

$$x_{j=1}|y = 3 \sim \text{Bernoulli}(\theta_{j=1,c=3}) \quad x_{j=2}|y = 3 \sim \text{Bernoulli}(\theta_{j=2,c=3})$$

$$\prod_{i=1}^5 \left(\left(\pi_{c=1}^{I(y_i=1)} \cdot p(x_{i,j=1}|\theta_{j=1,c=1}) \cdot p(x_{i,j=2}|\theta_{j=2,c=1}) \right) \cdot \left(\pi_{c=2}^{I(y_i=2)} \cdot p(x_{i,j=1}|\theta_{j=1,c=2}) \cdot p(x_{i,j=2}|\theta_{j=2,c=2}) \right) \cdot \left(\pi_{c=3}^{I(y_i=3)} \cdot p(x_{i,j=1}|\theta_{j=1,c=3}) \cdot p(x_{i,j=2}|\theta_{j=2,c=3}) \right) \right)$$

y	x_1	x_2
1	0	1
3	1	0
3	1	1
2	0	0
1	1	0

Let $m_c := \sum_{i=1}^m \mathbf{I}\{y^{(i)} = c\}$ be number of training examples in class c then,

$$\sum_{i=1}^m \log p(\mathcal{D}; \pi, \theta) = \sum_{c=1}^C m_c \log \pi_c + \sum_{c=1}^C \sum_{i: y^{(i)}=c} \sum_{d=1}^D \log p(x_d^{(i)}; \theta_{cd})$$

Log-likelihood function is concave in all parameters so...

1. Take derivatives with respect to π and θ separately.
2. Set derivatives to zero and solve

$$\hat{\pi}_c = \frac{m_c}{m}$$

**Fraction of training
examples from class c**

$$\hat{\theta}_{cd} = \frac{m_{cd}}{m_c}$$

**Number of “heads” in
training set from class c**

$$m_{cd} = \sum_{i=1}^m \mathbf{I}\{y^{(i)} = c, x_d^{(i)} = 1\}$$

Example: Naïve Bayes with Bernoulli Features

32

Analogy:

y	x_1	x_2	...	x_D		
5	0	1	1	0	1	0
3	1	0	0	0	0	0
3	1	1	1	1	0	0
...
4	0	0	1	1	0	1

$\hat{\pi}_c = \frac{m_c}{m}$
 $\hat{\theta}_{cd} = \frac{m_{cd}}{m_c}$
 $m_{cd} = \sum_{i=1}^m I\{y^{(i)} = c, x_d^{(i)} = 1\}$

$\hat{\pi}_3 = \frac{\# \text{ number of } y = 3}{\# \text{ total data}}$
 $\hat{\theta}_{3,2} = \frac{\# \text{ number of } y = 3 \text{ and } x_2 = 1}{\# \text{ number of } y = 3}$

Bernoulli Naïve Bayes: making prediction

$$\hat{\pi}_c = \frac{m_c}{m}$$

$$\hat{\theta}_{cd} = \frac{m_{cd}}{m_c}$$

Given one data point, it has 4 features (input), compare the probabilities:

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 0) &= p(y = 0) \cdot p(x_1, x_2, x_3, x_4 | y = 0) \\ &= p(y = 0) \cdot p(x_1 | y = 0) \cdot p(x_2 | y = 0) \cdot p(x_3 | y = 0) \cdot p(x_4 | y = 0) \end{aligned}$$

$$\begin{aligned} p(x_1, x_2, x_3, x_4, y = 1) &= p(y = 1) \cdot p(x_1, x_2, x_3, x_4 | y = 1) \\ &= p(y = 1) \cdot p(x_1 | y = 1) \cdot p(x_2 | y = 1) \cdot p(x_3 | y = 1) \cdot p(x_4 | y = 1) \end{aligned}$$

Scikit-learn has separate classes each feature type

`sklearn.naive_bayes.GaussianNB`

Real-valued features

`sklearn.naive_bayes.MultinomialNB`

Discrete K-valued feature counts (e.g. multiple die rolls)

`sklearn.naive_bayes.BernoulliNB`

Binary features (e.g. coinflip)

`sklearn.naive_bayes.CategoricalNB`

Discrete K-valued features (e.g. single die roll)

For large training data that don't fit in memory use Scikit-learn's out-of-core learning