



Computer
Science

CSC380: Principles of Data Science

Basic machine learning 2

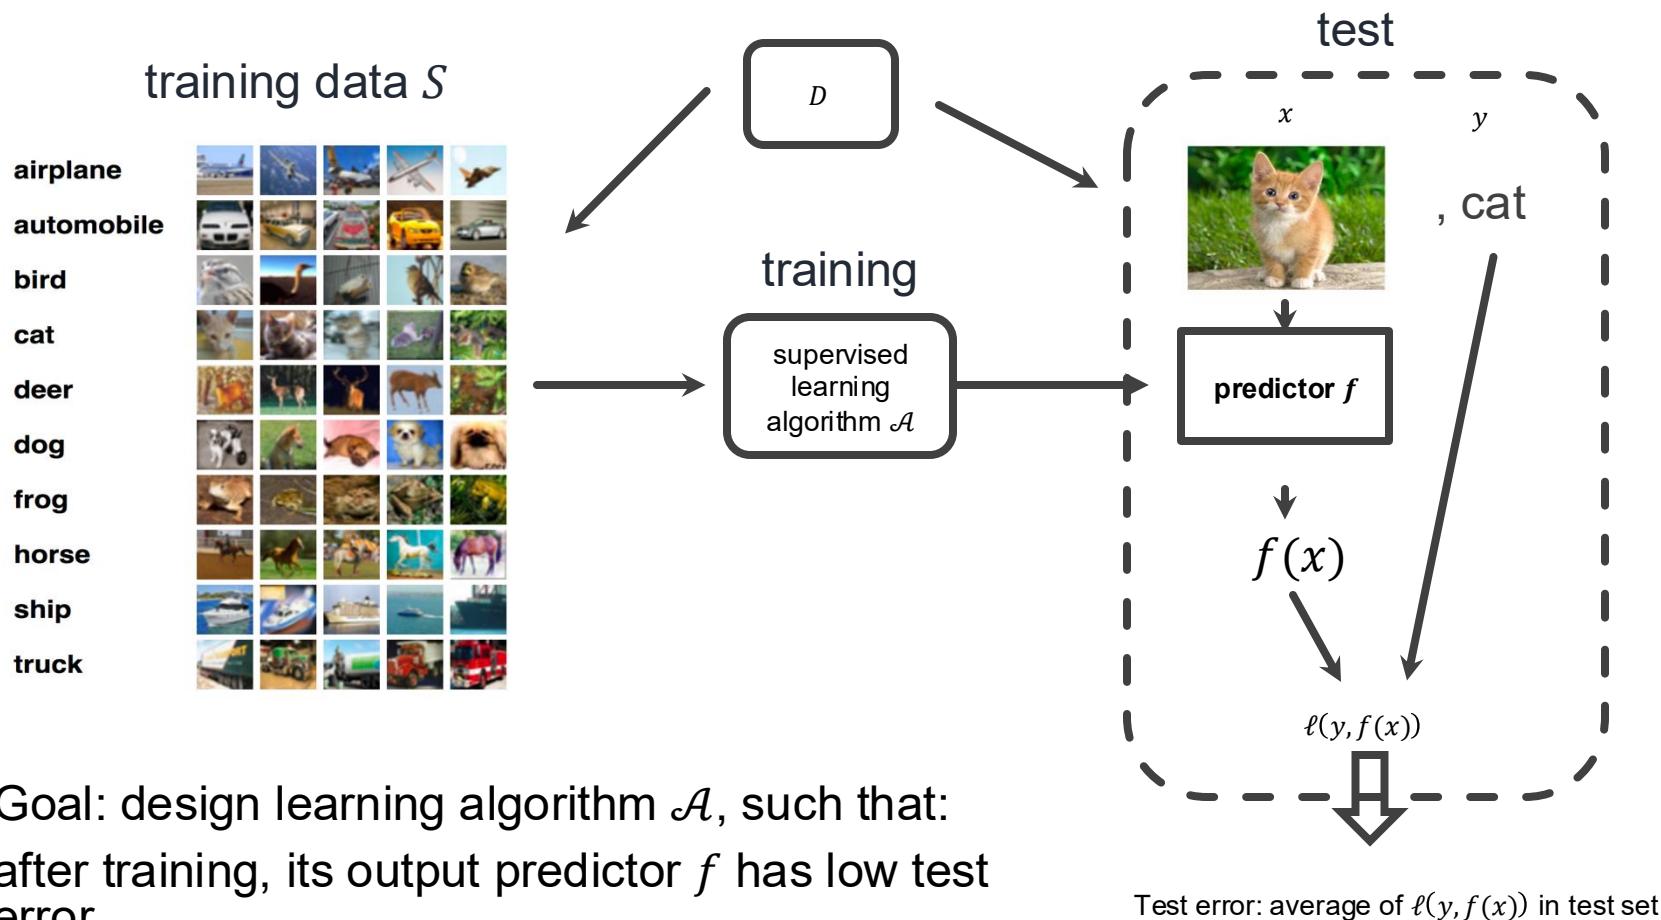
Xinchen Yu

- Classification basics
- Nearest neighbor Classification
- Logistic regression
- Classification: other considerations
 - Binary classification beyond accuracy
 - Multiclass classification

Classification recap

Supervised learning setup in one figure

4



Classification

- The labels are categorical
- Loss function ℓ : measures the quality of prediction \hat{y} respect to true label y
 - $\ell(y, \hat{y}) = I(y \neq \hat{y})$
 - I : indicator of predicate; 1 if true; 0 if false
- A classifier f 's error on a dataset S is the fraction of examples in S that it predicts incorrectly.
 - f 's training / test error is its error on training / test set
 - Accuracy = 1 – error

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



In-class activity: finding test error

A company develops a simple **spam classifier** f that predicts whether an email is **spam (1)** or **not spam (0)** based on the number of capital letters in the subject line.

f outputs **Spam** if the number of capital letters ≥ 5 , and **Not Spam** otherwise.

Suppose the test dataset is as follows. Find f 's test error.

Subject	True label	Predicted label
"WIN A FREE VACATION NOW!!!"	1	1
Meeting rescheduled to 3 PM	0	0
"HUGE DISCOUNT ON ALL ITEMS!!!"	1	1
URGENT: Please submit your report	0	1
Can you review this document?	0	0

$$f\text{'s test error} = 1/5 = 20\%$$

Nearest Neighbor Classification

	Rating	Easy?	AI?	Sys?	Thy?	Morning?
Label: "like"	+2	y	y	n	y	n
	+2	y	y	n	y	n
	+2	n	y	n	n	n
	+2	n	n	n	y	n
	+2	n	y	y	n	y
	+1	y	y	n	n	n
	+1	y	y	n	y	n
	+1	n	y	n	y	n
	0	n	n	n	n	y
	0	y	n	n	y	y
Label: "dislike"	0	n	y	n	y	n
	0	y	y	y	y	y
	-1	y	y	y	n	y
	-1	n	n	y	y	n
	-1	n	n	y	n	y
	-1	y	n	y	n	y
	-2	n	n	y	y	n
	-2	n	y	y	n	y
	-2	y	n	y	n	n
	-2	y	n	y	n	y

Features

Suppose we'd like to build a recommendation system for classes

We've collected information about many past classes

We can frame this as a classification problem:

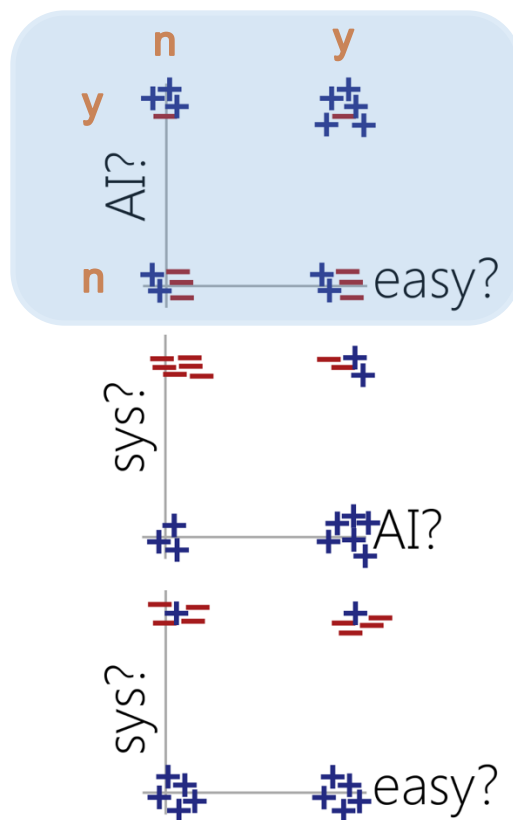
Predict like/dislike from class features

Example: Course Recommendation

9

	Rating	Easy?	AI?	Sys?	Thy?	Morning?
Label: +	+2	y	y	n	y	n
	+2	y	y	n	y	n
	+2	n	y	n	n	n
	+2	n	n	n	y	n
	+2	n	y	y	n	y
	+1	y	y	n	n	n
	+1	y	y	n	y	n
	+1	n	y	n	y	n
	0	n	n	n	n	y
	0	y	n	n	y	y
Label: -	0	n	y	n	y	n
	0	y	y	y	y	y
	-1	y	y	y	n	y
	-1	n	n	y	y	n
	-1	n	n	y	n	y
	-1	y	n	y	n	y
	-2	n	n	y	y	n
	-2	n	y	y	n	y
	-2	y	n	y	n	n
	-2	y	n	y	n	y

Features




Each course's feature is Represented as points in 5-dimensional space

That's too many dimensions to plot...so we look at 2D projections...

Observation: examples with same labels tend to be closer!

Nearest neighbor classification

- Given a new course, would like to predict its label (+/-)
- Idea: Find its most similar course in the training set, and use that course's label to predict




- Oftentimes convenient to work with feature $x \in \mathbb{R}^d$
- Distances in \mathbb{R}^d :


notation $x(f)$: $x = (x(1), \dots, x(d))$

- (popular) Euclidean distance $d_2(x, x') = \sqrt{\sum_{f=1}^d (x(f) - x'(f))^2}$
- Manhattan distance $d_1(x, x') = \sum_{f=1}^d |x(f) - x'(f)|$

- How to extract features as **real values**?

- Boolean features: $\{Y, N\} \rightarrow \{0, 1\}$
- Categorical features: $\{\text{Red, Blue, Green, Black}\}$
 - Convert to $\{1, 2, 3, 4\}$?
 - Better one-hot encoding: $(1, 0, 0, 0), \dots, (0, 0, 0, 1)$
(IsRed?/isGreen?/isBlue?/IsBlack?)






Q: Can we predict using 1 nearest neighbor's?


Query point ? Will be classified as + but should be -

Problem: predicting using 1 nearest neighbor's label can be sensitive to noisy data

How to mitigate this?

- Training set: $S = \{ (x_1, y_1), \dots, (x_m, y_m) \}$
- **Key insight:** given test example x , its label should resemble the labels of *nearby points*
- Function
 - input: x
 - find the k nearest points to x from S ; call their indices $N(x)$
 - output:
 - (classification) the majority vote of $\{y_i: i \in N(x)\}$
 - (regression) the average of $\{y_i: i \in N(x)\}$








decision boundary

- Features having different scales can be problematic.

- Ex: ski vs. snowboard classification

$$d = \sqrt{(height_1 - height_2)^2 + (weight_1 - weight_2)^2}$$



- One solution: feature standardization

- Features having different scale can be problematic


- [Definition] **Standardization**

- For each feature f , compute $\mu_f = \frac{1}{m} \sum_{i=1}^m x_f^{(i)}$, $\sigma_f = \sqrt{\frac{1}{m} \sum_{i=1}^m (x_f^{(i)} - \mu_f)^2}$
- Then, transform the data by $\forall f \in \{1, \dots, d\}, \forall i \in \{1, \dots, m\}, x_f^{(i)} \leftarrow \frac{x_f^{(i)} - \mu_f}{\sigma_f}$


after transformation, each feature has mean 0 and variance 1

- Be sure to keep the “standardize” function and apply it to the test points.
 - Save $\{(\mu_f, \sigma_f)\}_{f=1}^d$
 - For test point x^* , apply $x_f^* \leftarrow \frac{x_f^* - \mu_f}{\sigma_f}, \forall f$

here's a case in which there is one relevant feature x_1 and a 1-NN rule classifies each instance correctly




consider the effect of an irrelevant feature x_2 on distances and nearest neighbors




- Mitigation: feature selection

- Q: How would a k -NN classifier predict when k =training set size?
 - Predict majority label everywhere
 - Underfitting
- Q: What is the training error of a 1-NN classifier?
 - 0
 - Overfitting



k can be viewed as a model complexity measure

Smaller k results in a more complex model




Issue 3: choosing k

We'd like to choose appropriate k to balance model bias and complexity

We can choose k in the same way we chose λ in ridge regression

- Cross validation



Scikit-learn nearest neighbors

```
class sklearn.neighbors.NearestNeighbors(*, n_neighbors=5, radius=1.0,  
algorithm='auto', leaf_size=30, metric='minkowski', p=2, metric_params=None,  
n_jobs=None)
```

[\[source\]](#)

Unsupervised learner for implementing neighbor searches.

```
# 1. Load the Iris dataset  
iris = load_iris()  
X = iris.data # Features  
y = iris.target # Target labels (species)  
  
# 2. Split the dataset into training and testing sets (80% train, 20% test)  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# 3. Create the KNN classifier model  
knn = KNeighborsClassifier(n_neighbors=3) # Use 3 nearest neighbors  
  
# 4. Train the model on the training data  
knn.fit(X_train, y_train)
```



Scikit-learn nearest neighbors

```
# 5. Make predictions on the test set
y_pred = knn.predict(X_test)

# 6. Evaluate the model using accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy of the KNN model: {accuracy * 100:.2f}%')

# Optionally, display the predictions vs. actual values
print(f'Predictions: {y_pred}')
print(f'Actual: {y_test}')
```

Accuracy of the KNN model: 100.00%

Predictions: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 2 0 0]

Actual: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 2 0 0]