Computer Science

CSC380: Principles of Data Science

Basic machine learning 2

Xinchen Yu

Outline

- Classification basics
- Nearest neighbor Classification
- Logistic regression

Classification recap

Supervised learning setup in one figure

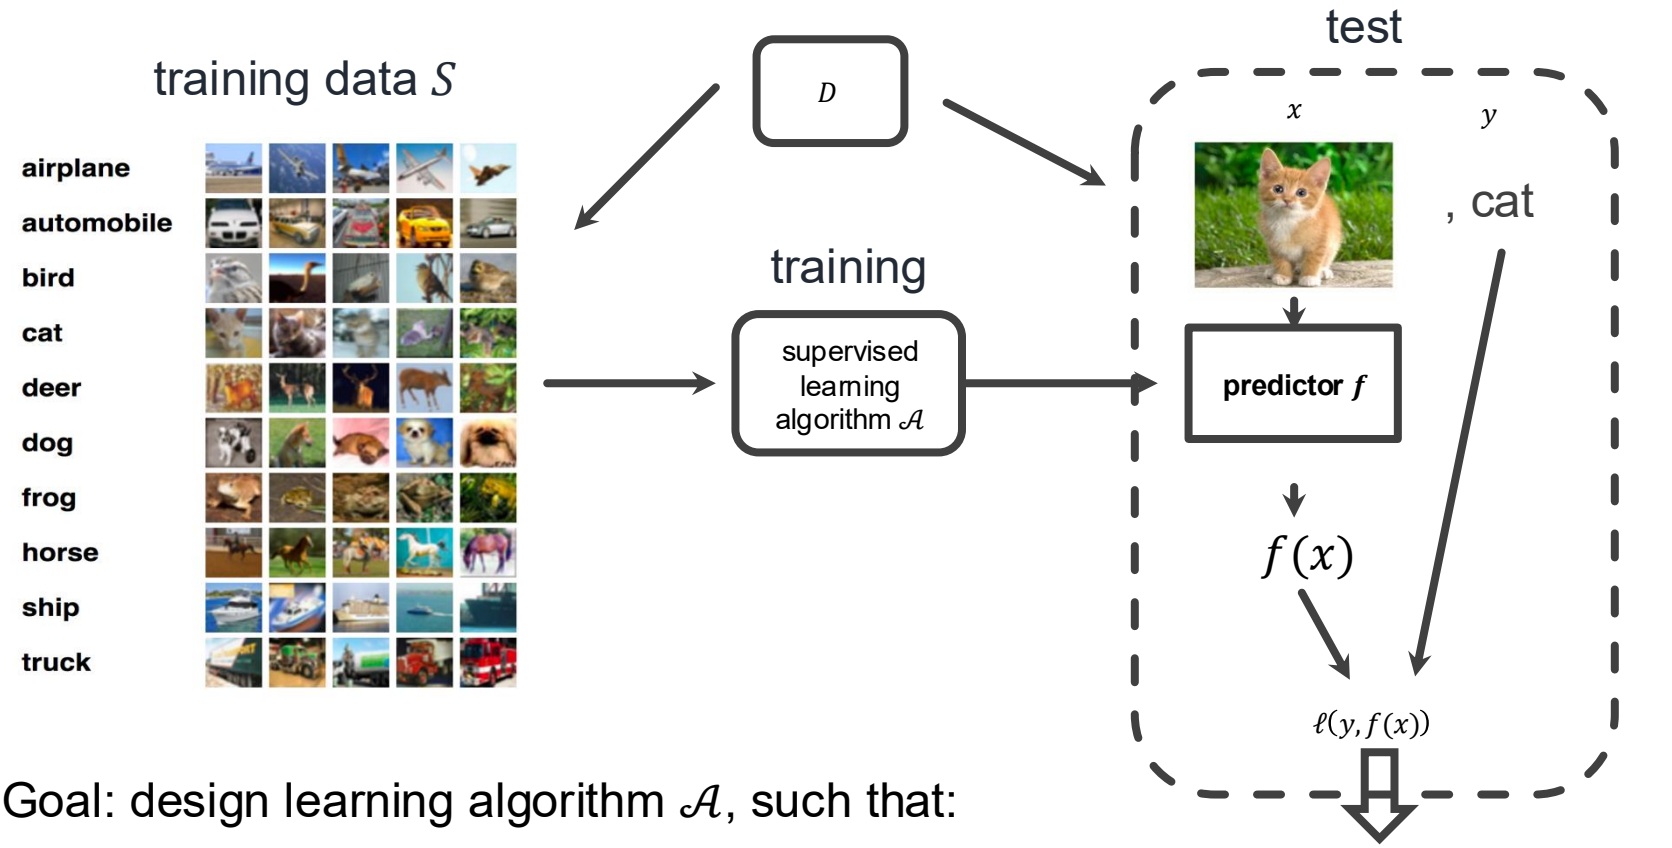

- Goal: design learning algorithm $\mathcal{A}$, such that:
- after training, its output predictor f has low test error

Test error: average of $\ell(y, f(x))$ in test set

Classification

- The labels are categorical
- Loss function ℓ: measures the quality of prediction $\hat{y}$ respect to true label y
  - $\ell(y, \hat{y}) = I(y \neq \hat{y})$
  - I: indicator of predicate; 1 if true; 0 if false

- A classifier f's error on a dataset S is the fraction of examples in S that it predicts incorrectly.
  - f's training / test error is its error on training / test set
  - Accuracy = 1 – error

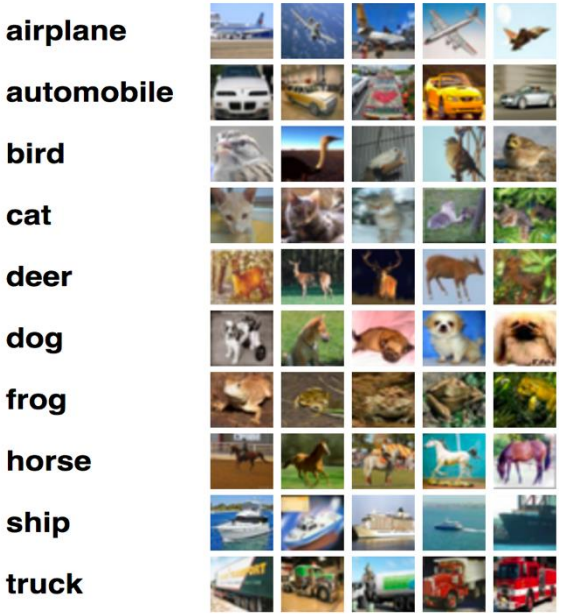

In-class activity: finding test error

A company develops a simple **spam classifier** f that predicts whether an email is **spam (1)** or **not spam (0)** based on the number of capital letters in the subject line.

f outputs **Spam** if the number of capital letters ≥ 5, and **Not Spam** otherwise.

Suppose the test dataset is as follows. Find f's test error.

Subject	True label	Predicted label
"WIN A FREE VACATION NOW!!!"	1	1
Meeting rescheduled to 3 PM	0	0
"HUGE DISCOUNT ON ALL ITEMS!!!"	1	1
URGENT: Please submit your report	0	1
Can you review this document?	0	0

f's test error = 1/5 = 20%

Nearest Neighbor Classification

Example: Course Recommendation

Rating	Easy?	AI?	Sys?	Thy?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y

Label: "like"

Label: "dislike"

Features

Suppose we'd like to build a recommendation system for classes

We've collected information about many past classes

We can frame this as a classification problem:

Predict like/dislike from class features

Example: Course Recommendation

Rating	Easy?	AI?	Sys?	Thy?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y

Label: "like"

Label: "dislike"

Features

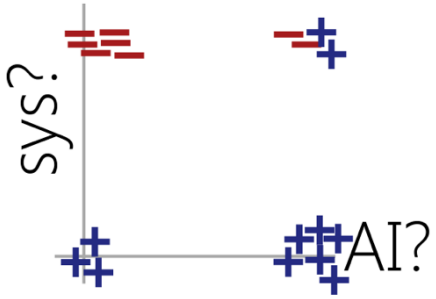

Each feature is represented as points in 5-dimensional space

We only look at 2D projections…

Observation: examples with same labels tend to be closer!

Example: Course Recommendation

	Rating	Easy?	AI?	Sys?	Thy?	Morning?
Label: +	+2	y	y	n	y	n
	+2	y	y	n	y	n
	+2	n	y	n	n	n
	+2	n	n	n	y	n
	+2	n	y	y	n	y
	+1	y	y	n	n	n
	+1	y	y	n	y	n
	+1	n	y	n	y	n
	0	n	n	n	n	y
	0	y	n	n	y	y
	0	n	y	n	y	n
	0	y	y	y	y	y
Label: -	-1	y	y	y	n	y
	-1	n	n	y	y	n
	-1	n	n	y	n	y
	-1	y	n	y	n	y
	-2	n	n	y	y	n
	-2	n	y	y	n	y
	-2	y	n	y	n	n
	-2	y	n	y	n	y

Features

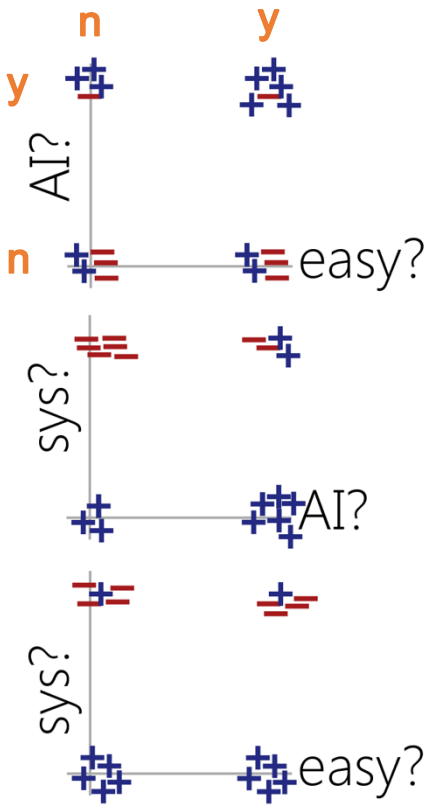

Each course's feature is Represented as points in 5-dimensional space

That's too many dimensions to plot…so we look at 2D projections…

Observation: examples with same labels tend to be closer!

Nearest neighbor classification

- Given a new course, would like to predict its label (+/-)
- Idea: Find its most similar course in the training set, and use that course's label to predict

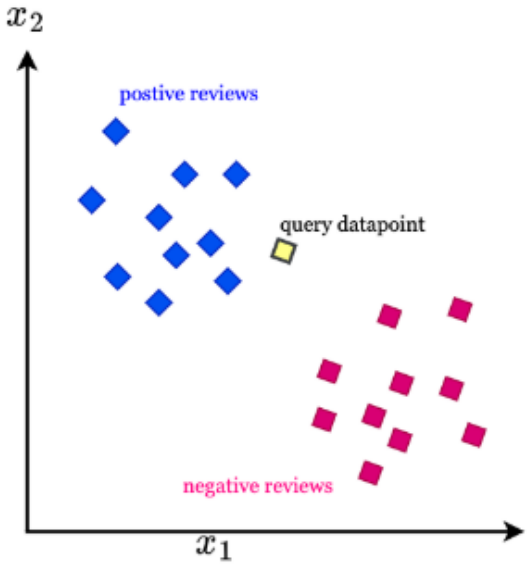

Measuring nearest neighbors

- Oftentimes convenient to work with feature $x \in \mathrm{R}^d$
- Distances in R^d:

notation $x(f)$: $x = (x(1), \ldots, x(d))$

  - (popular) Euclidean distance $d_2(x, x') = \sqrt{\sum_{f=1}^{d} \left(x(f) - x'(f)\right)^2}$
  - Manhattan distance $d_1(x, x') = \sum_{f=1}^{d} |x(f) - x'(f)|$

- How to extract features as **<u>real values</u>**?
  - Boolean features: {Y, N} -> {0,1}
  - Categorical features: {Red, Blue, Green, Black}
    - one-hot encoding: (1,0,0,0), .., (0,0,0,1)

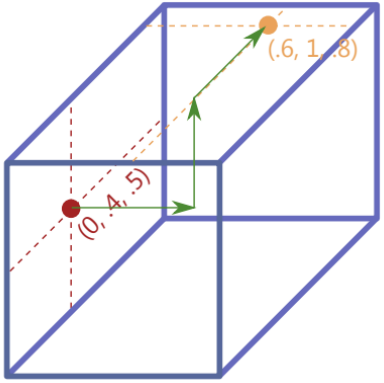

k-nearest neighbors (k-NN): main concept

- Training set: $S = \{ (x_1, y_1), \ldots, (x_m, y_m) \}$

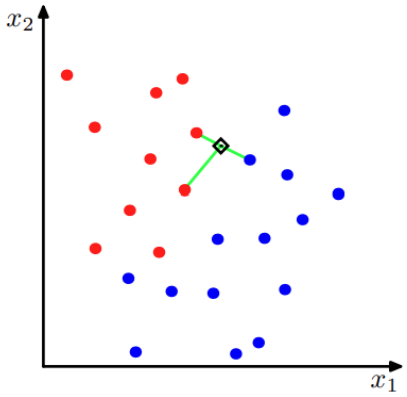

- **Key insight**: given test example x, its label should resemble the labels of *nearby points*

- Function
  - input: x
  - find the k nearest points to x from S; call their indices $N(x)$
  - output:
    - (classification) the **majority vote** of $\{y_i : i \in N(x)\}$
    - (regression) the average of $\{y_i : i \in N(x)\}$

k-NN classification example

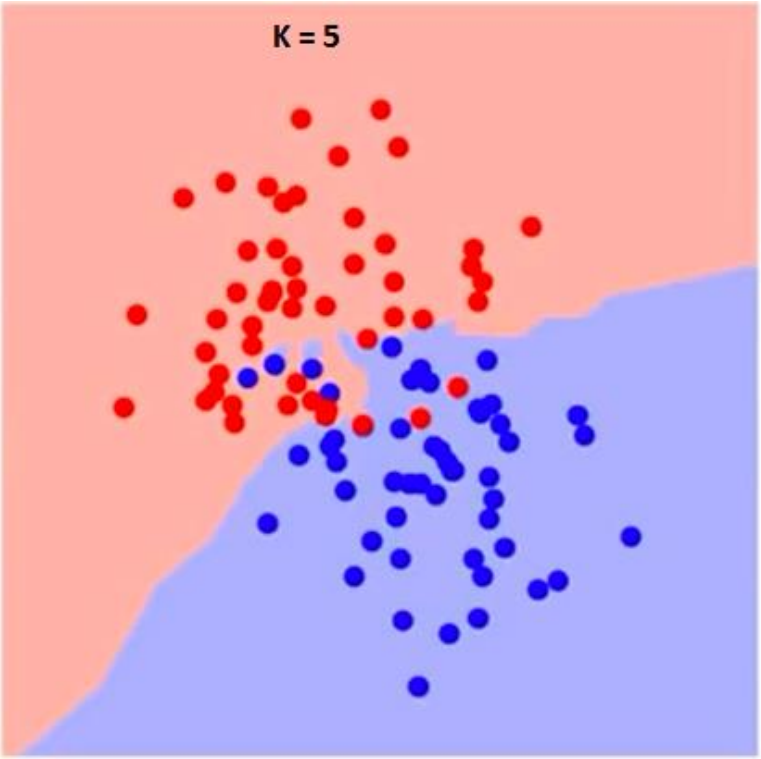

decision boundary

Issue 1: scaling

- Features having different scales can be problematic.

- Ex: ski vs. snowboard classification

$$d = \sqrt{(height_1 - height_2)^2 + (width_1 - width_2)^2}$$

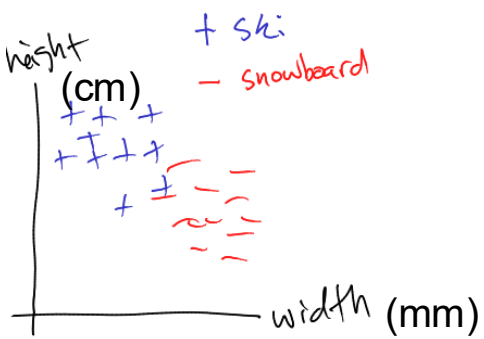

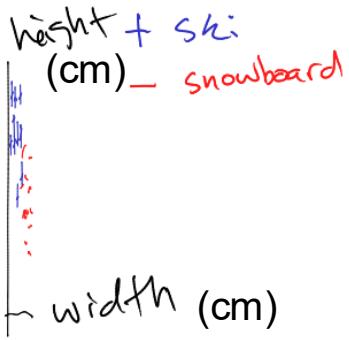

- One solution: feature standardization

Make sure features are scaled fairly

- **Standardization**
  - For each feature f, compute $\mu_f = \frac{1}{m}\sum_{i=1}^{m} x_f^{(i)}$, $\sigma_f = \sqrt{\frac{1}{m}\sum_{i=1}^{m}\left(x_f^{(i)} - \mu_f\right)^2}$
  - Then, transform the data by $\forall f \in \{1, \dots, d\}, \forall i \in \{1, \dots, m\}$, $x_f^{(i)} \leftarrow \frac{x_f^{(i)} - \mu_f}{\sigma_f}$

after transformation, each feature has mean 0 and variance 1

- Be sure to keep the "standardize" function and apply it to the test points.
  - Save $\{(\mu_f, \sigma_f)\}_{f=1}^{d}$
  - For test point x^*, apply $x_f^* \leftarrow \frac{x_f^* - \mu_f}{\sigma_f}, \forall f$

Issue 2: irrelevant features

here's a case in which there is one relevant feature x_1 and a 1-NN rule classifies each instance correctly

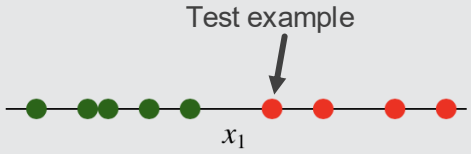

consider the effect of an irrelevant feature x_2 on distances and nearest neighbors

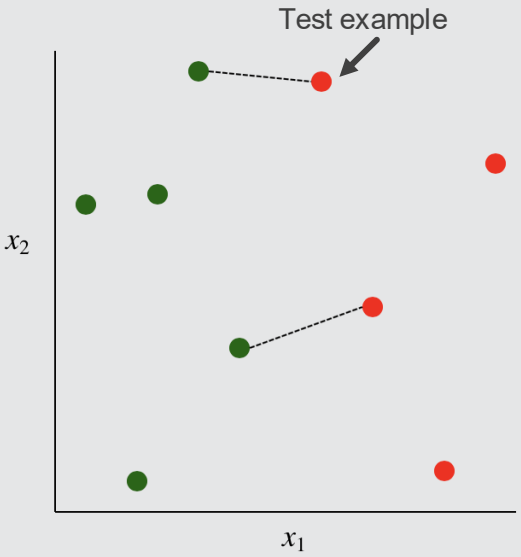

- Mitigation: feature selection

Issue 3: choosing k

- Q: How would a k-NN classifier predict when k=training set size?
  - Predict majority label everywhere
  - Underfitting

- Q: What is the training error of a 1-NN classifier?
  - 0
  - Overfitting

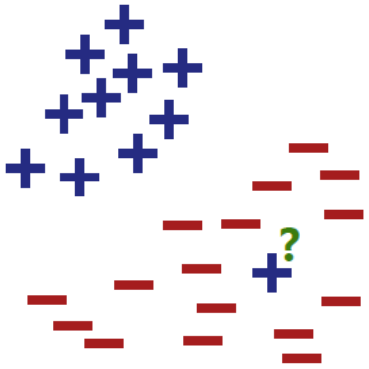

Issue 3: choosing k

k can be viewed as a model complexity measure

Smaller k results in a more complex model

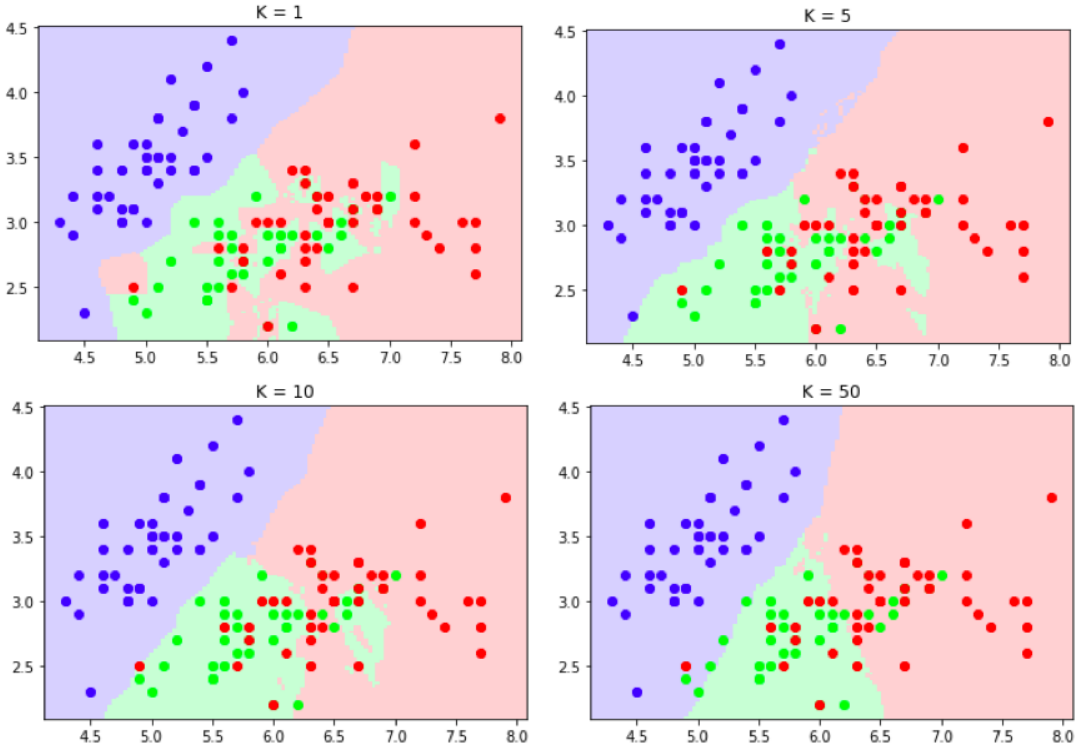

Issue 3: choosing k

We'd like to choose appropriate k to balance model bias and complexity

We can choose k in the same way we chose λ in ridge regression

- Cross validation

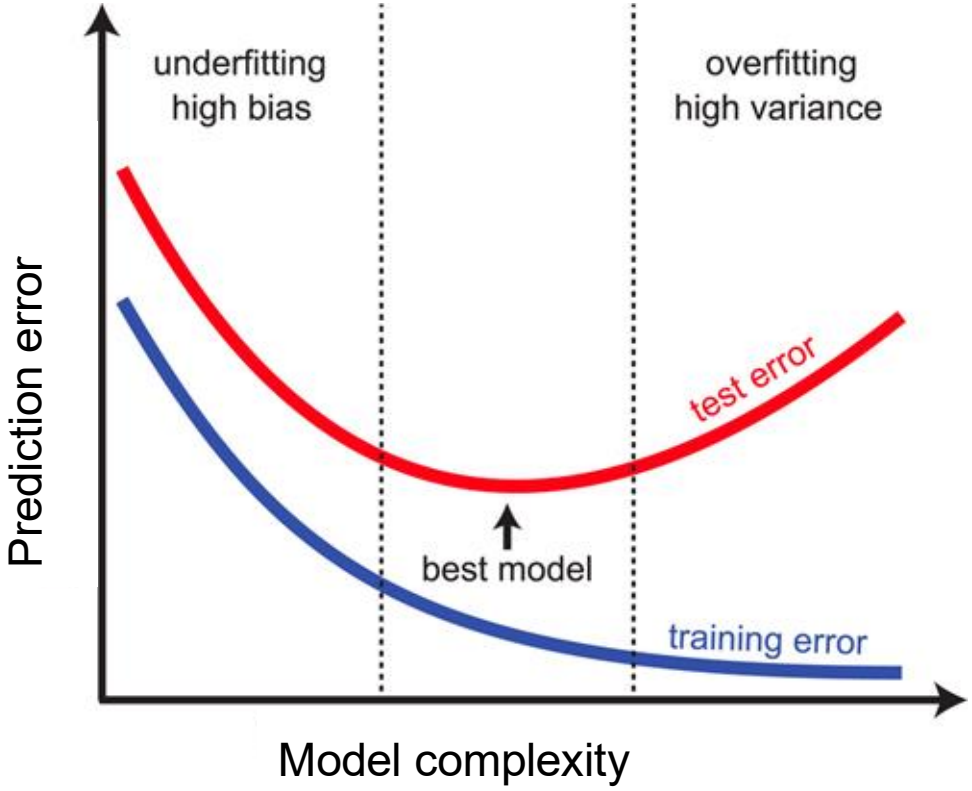

Scikit-learn nearest neighbors

class sklearn.neighbors.**NearestNeighbors**(*, *n_neighbors*=5, *radius*=1.0, *algorithm*='auto', *leaf_size*=30, *metric*='minkowski', *p*=2, *metric_params*=None, *n_jobs*=None) [source]

Unsupervised learner for implementing neighbor searches.

```
# 1. Load the Iris dataset
iris = load_iris()
X = iris.data  # Features
y = iris.target  # Target labels (species)

# 2. Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 3. Create the KNN classifier model
knn = KNeighborsClassifier(n_neighbors=3)  # Use 3 nearest neighbors

# 4. Train the model on the training data
knn.fit(X_train, y_train)
```

Scikit-learn nearest neighbors

```python
# 5. Make predictions on the test set
y_pred = knn.predict(X_test)

# 6. Evaluate the model using accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy of the KNN model: {accuracy * 100:.2f}%')

# Optionally, display the predictions vs. actual values
print(f'Predictions: {y_pred}')
print(f'Actual: {y_test}')
```

```
Accuracy of the KNN model: 100.00%
Predictions: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 2 0 0]
Actual: [1 0 2 1 1 0 1 2 1 1 2 0 0 0 0 1 2 1 1 2 0 2 0 2 2 2 2 2 0 0]
```

Logistic regression

Classification with logistic regression

Training data: number of hours studied for the course.

Labels: Pass (1) or Fail (0)

Task: Pass (1) or Fail (0)?

Classification with logistic regression

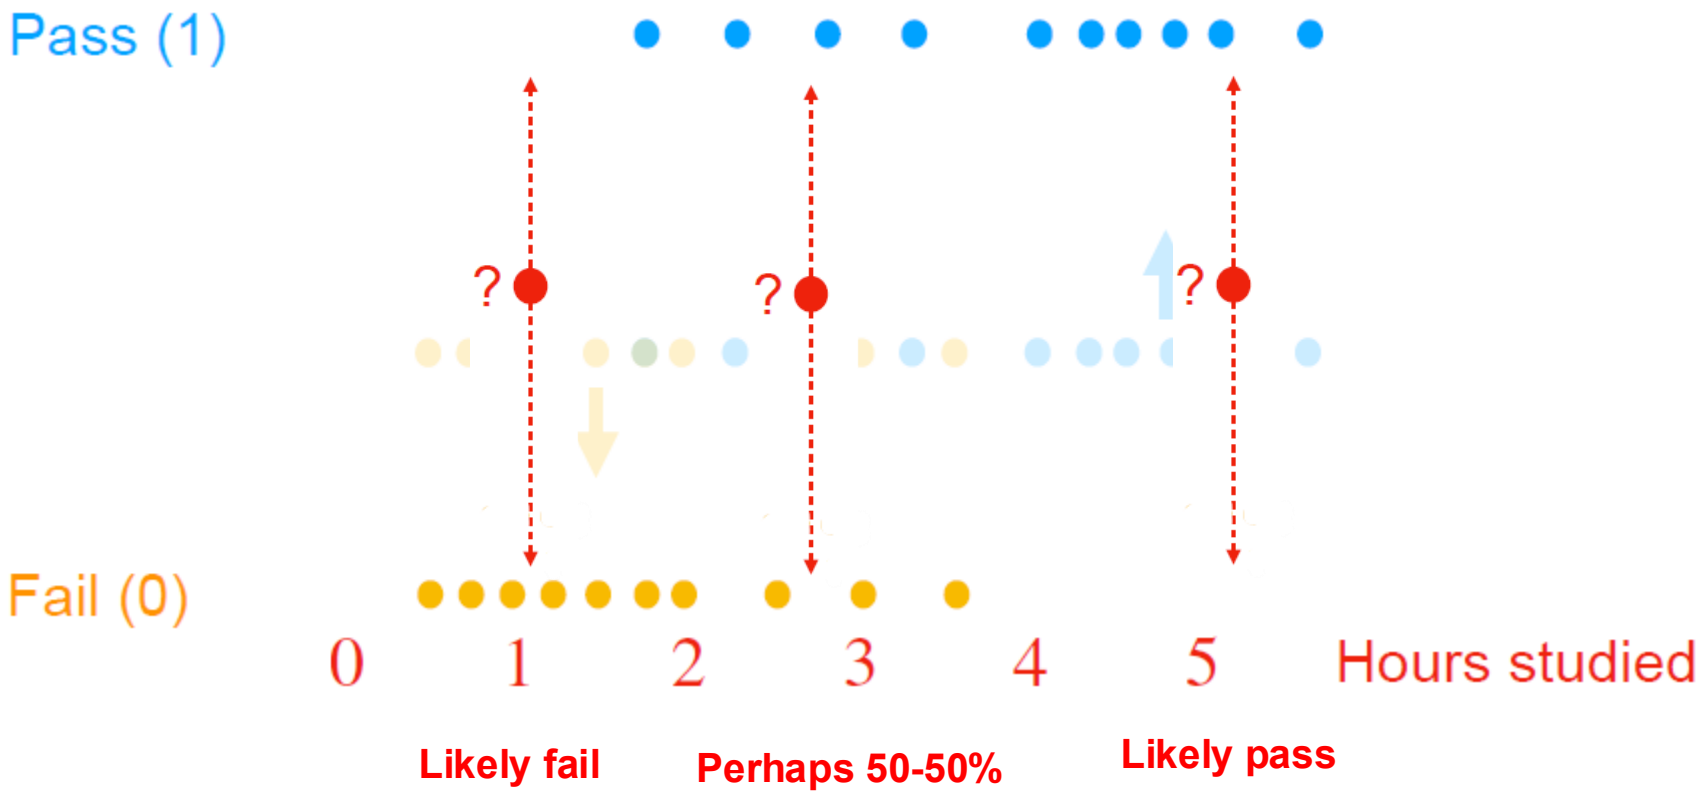

Classification with logistic regression

$Y = 1$

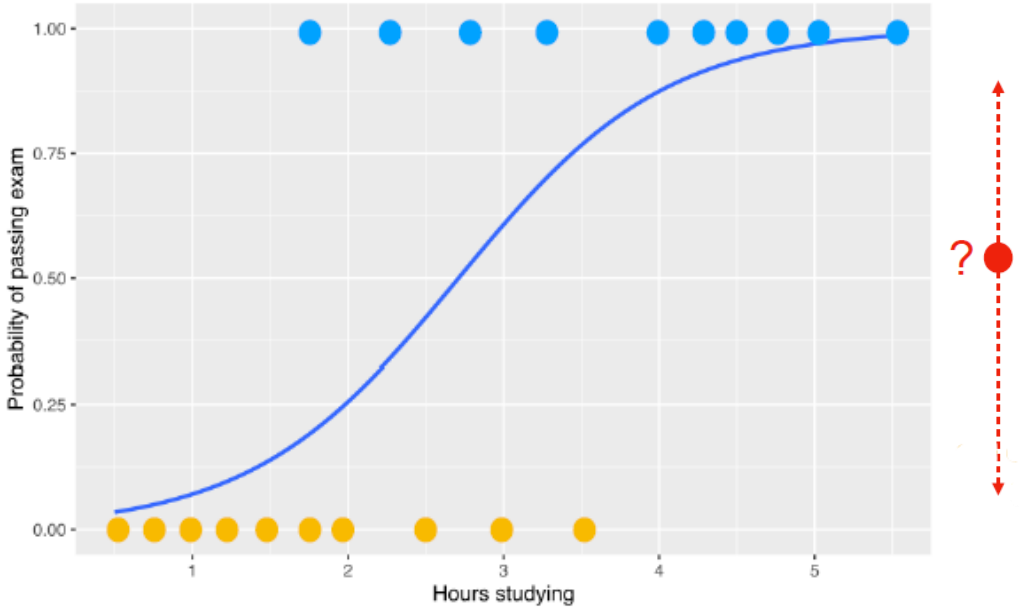

$Y = 0$

X: feature

Y: label

Blue curve:
$P(Y = 1 \mid X = x)$

Prediction:

If prob < 0.5: fail

Else: pass

What is a reasonable form of $P(Y = 1 \mid X = x)$?

Classification with logistic regression

We will assume that:

$$P(Y = 1 \mid X = x) = \sigma(w \cdot x + b) = \frac{1}{1 + e^{-(w \cdot x + b)}}$$

i.e., $\sigma(w \cdot x + b)$, for some w, b

For d-dim x, this is dot product

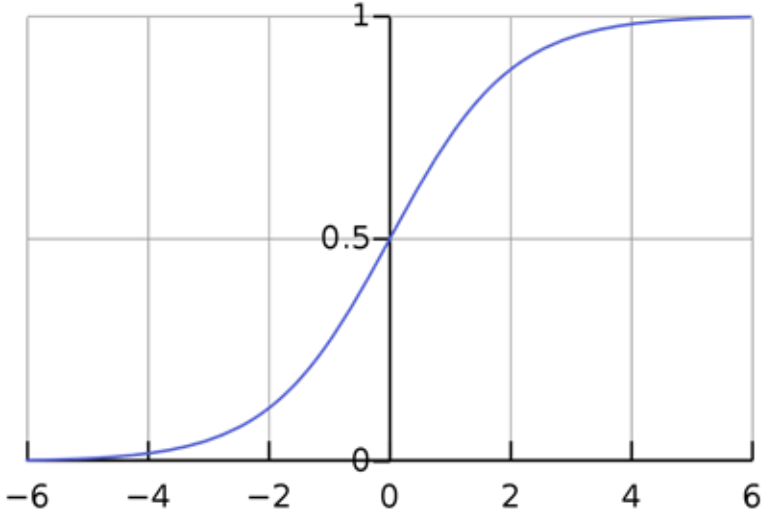

$\sigma(z) := \frac{1}{1+e^{-z}}$ is the *logistic function*

Logistic regression model

w controls the shape of the probability curve

$$p = \frac{1}{1 + e^{-10\,x}}$$

$$p = \frac{1}{1 + e^{-3\,x}}$$

$$p = \frac{1}{1 + e^{10\,x}}$$

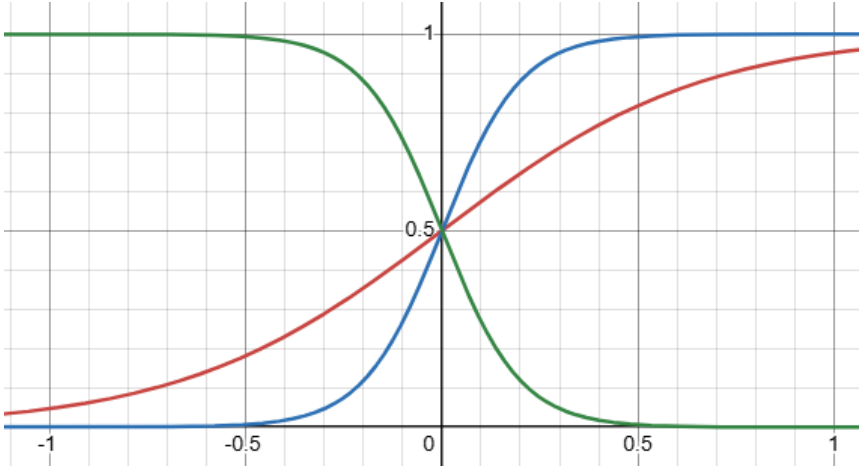

b controls the location of the probability curve

$$p = \frac{1}{1 + e^{-10\,x}}$$

$$p = \frac{1}{1 + e^{-10\,(x+1)}}$$

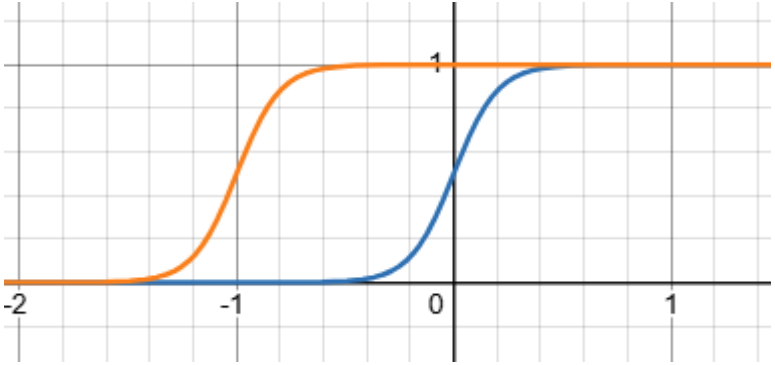

Classification with logistic regression

Example Suppose we fit logistic regression model with b = 0.15 and w = 0.575. What is the model's predicted probability that a student who have studied for x = 2 hours passes?

$P(Y = 1 \mid X = x) = \frac{1}{1+e^{-z}}$, where $z = w \cdot x + b = 1$

Thus, the predicted pass prob = $\frac{1}{1+e^{-1}} = 0.73$

Fitting a logistic regression model

- Recall: loss for linear regression was $\frac{1}{n}\sum_i (y_i - w \cdot x_i)^2$

- How about logistic regression?
  - y_i's are in 0, 1

$Y = 1$

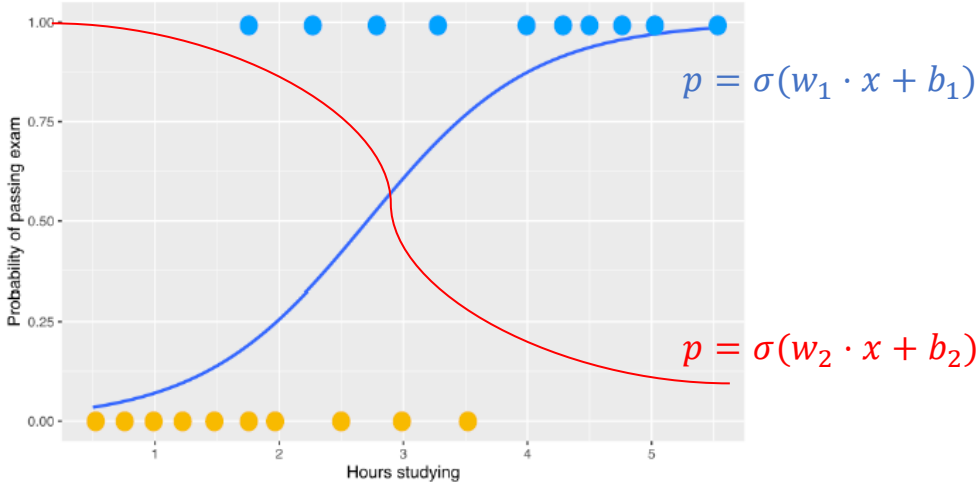

$Y = 0$

Which logistic regression model fits data better, red or blue?

Fitting a logistic regression model

We'd like to choose w and b such that:

for x whose label is more likely to be 1 (Y = 1)

=> p is large (blue): <u>penalize less</u>; p is small (red): penalize more

for x whose label is more likely to be 0 (Y = 0)

=> p is large (red): penalize more; p is small (blue): <u>penalize less</u>

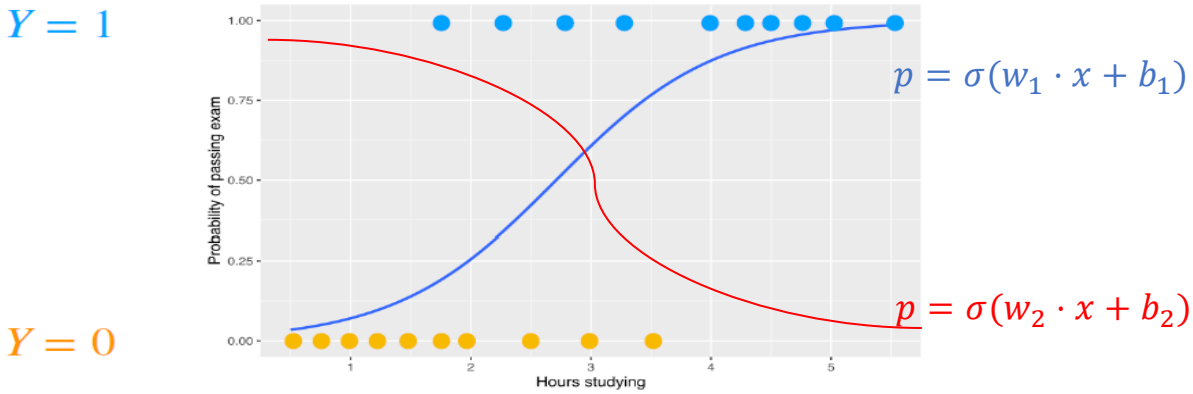

Need a loss function that penalizes predictions differently depending on the true label

Cross entropy loss

$$\ell(y,p) = y \ln\frac{1}{p} + (1-y)\ln\frac{1}{1-p}$$

$$= \begin{cases} \ln\frac{1}{p}, & y = 1 \\ \ln\frac{1}{1-p}, & y = 0 \end{cases}$$

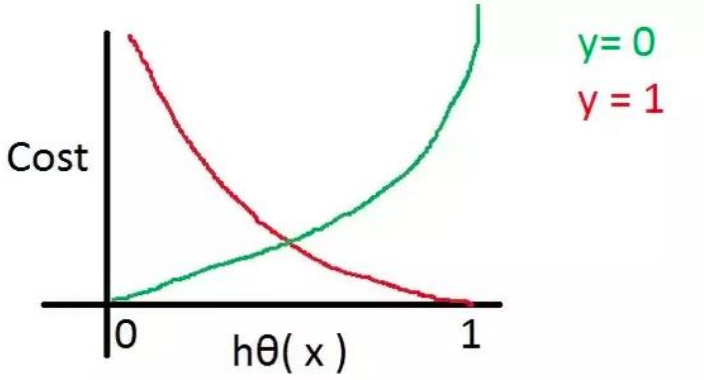

$p_i = \sigma(w \cdot x_i + b)$

Minimizing CE loss incentivizes the model's predictive probability to align with labels

Fitting a logistic regression model

- We find w and b to minimize:

$$\sum_i \left(y_i \ln \frac{1}{p_i} + (1 - y_i) \ln \frac{1}{1-p_i} \right),$$

Logistic loss, aka Cross-entropy (CE) loss

where $p_i = P(Y = 1 \mid x_i) = \frac{1}{1+e^{-(w \cdot x_i + b)}}$

- What is the i-th example's loss when:
  - $y_i = 1$ and $p_i \approx 1$? ≈ 0
  - $y_i = 1$ and $p_i \approx 0$? Large
  - $y_i = 0$ and $p_i \approx 1$? Large
  - $y_i = 0$ and $p_i \approx 0$? ≈ 0

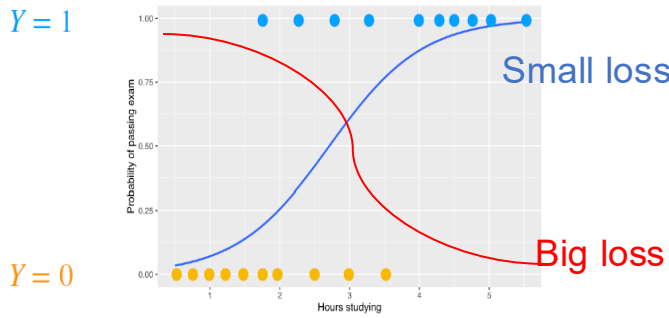

sklearn.linear_model.LogisticRegression

class sklearn.linear_model.LogisticRegression(*penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None*) [source]

penalty : *{'l1', 'l2', 'elasticnet', 'none'}, default='l2'*

Specify the norm of the penalty:

- `'none'`: no penalty is added;
- `'l2'`: add a L2 penalty term and it is the default choice;
- `'l1'`: add a L1 penalty term;
- `'elasticnet'`: both L1 and L2 penalty terms are added.

Similar to linear regression, add regularization to combat overfitting

$$\text{argmin}_w \text{ Logistic} - \text{Loss}(w) + \lambda|w|$$

tol : *float, default=1e-4*

Tolerance for stopping criteria.

C : *float, default=1.0*

$C = 1/\lambda$

Inverse of regularization strength; must be a positive float. Like in support vector machines, smaller values specify stronger regularization.

Scikit-Learn Logistic Regression

log_regression = sklearn.linear_model.LogisticRegression()

```
_ = log_regression.fit(pd.DataFrame(x), y)

y_pred = log_regression.predict_proba(pd.DataFrame(x))
log_y_pred_1 = [item[1] for item in y_pred]

fig = plt.figure(figsize=(10,5))
xlabel = 'Age'
ylabel = 'Purchased'
plt.xlabel(xlabel)
plt.ylabel(ylabel)
plt.grid(color='k', linestyle=':', linewidth=1)
plt.plot(x, y, 'xb')
plt.plot(x, log_y_pred_1, '-r')
_ = plt.plot(x, line_point_5,'-g')
```

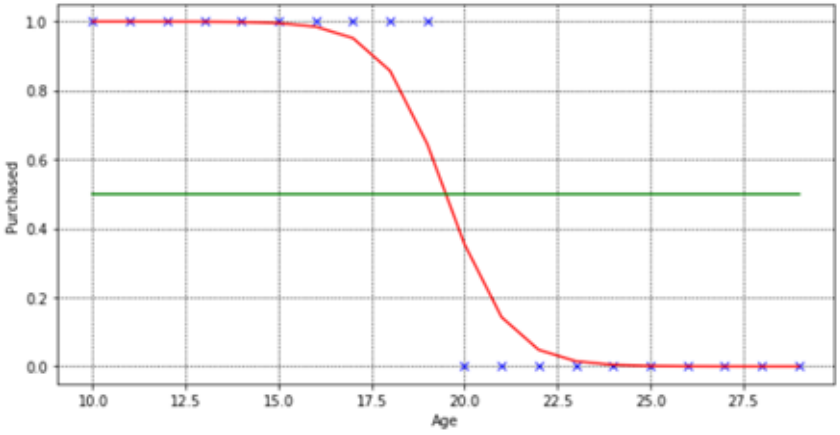

Function `predict_proba(X)` returns prediction of class assignment probabilities for each class. It returns n by C matrix if n data points were provided as argument.

(C=number classes)

https://towardsdatascience.com/why-linear-regression-is-not-suitable-for-binary-classification-c64457be8e28